



logstash



elasticsearch.



Microservices

By Narendranath Reddy


Name: Narendranath Reddy Thota

- ❖ d3 Whitepaper Core Team
- ❖ Trainer from Gama
- ❖ Blockchain Full Stack Developer
- ❖ Technology Analyst
- ❖ Software Developer
- ❖ Technology Speaker

MAERSK

Email: narendranath.thota@ust-global.com

P: +917288838869




Multi-Language | Run Anywhere | In-Memory



“The Monolith”

Q. Clipboard



Challenges with monolithic software

Difficult to
scale

Architecture is
hard to maintain
and evolve

Lack of agility

Long
Build/Test/Release
Cycles
(who broke the build?)

New releases
take months

Lack of innovation

Operations
is a nightmare
(module X is failing,
who's the owner?)

Long time to add
new features

Frustrated customers

Monolith development lifecycle

developers




delivery pipeline

build


test

release

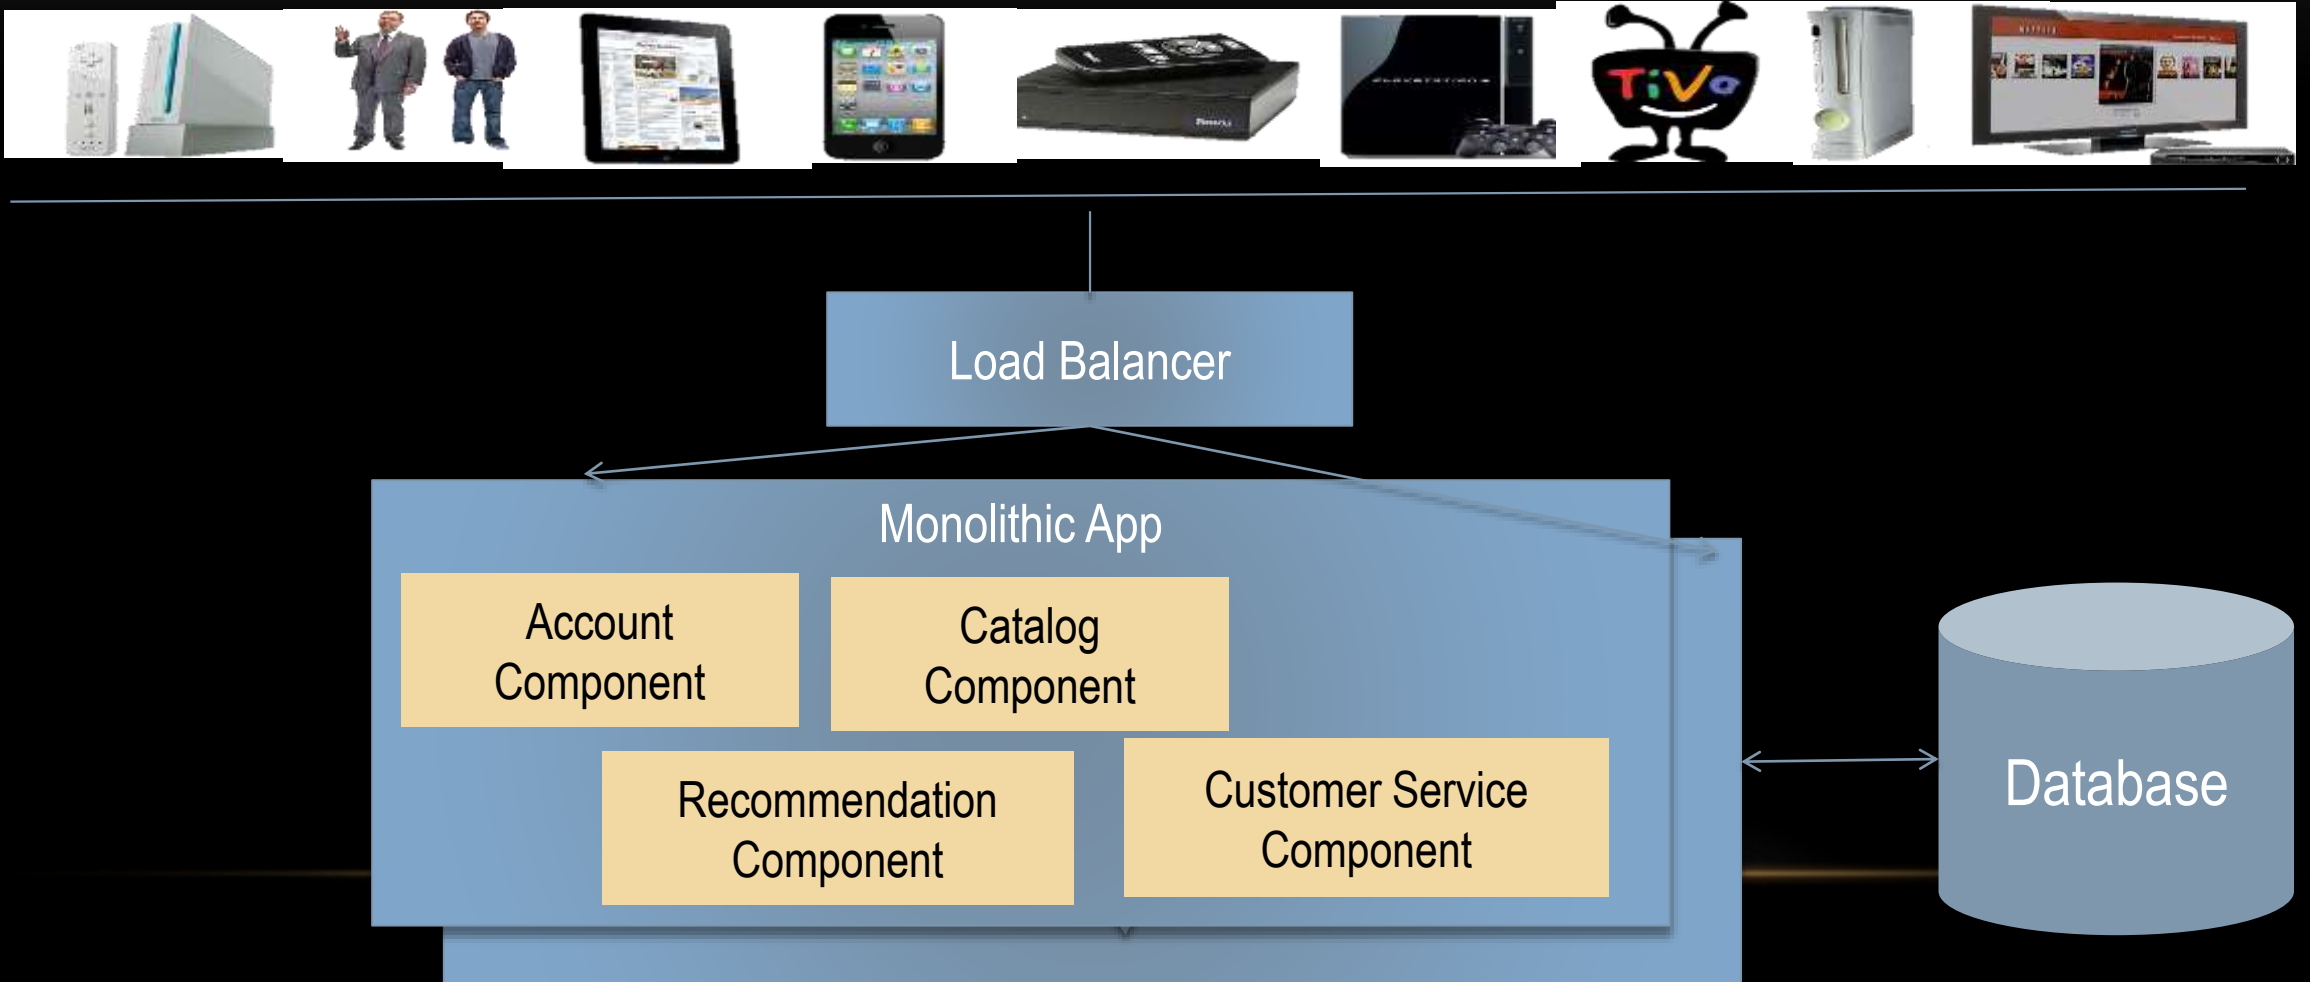
app
(aka the "monolith")



Too much software coupling



Monolithic Architecture



Comparing Monolithic to MicroServices




MONOLITHIC APP (VARIOUS COMPONENTS LINKED TOGETHER)



CHANGE

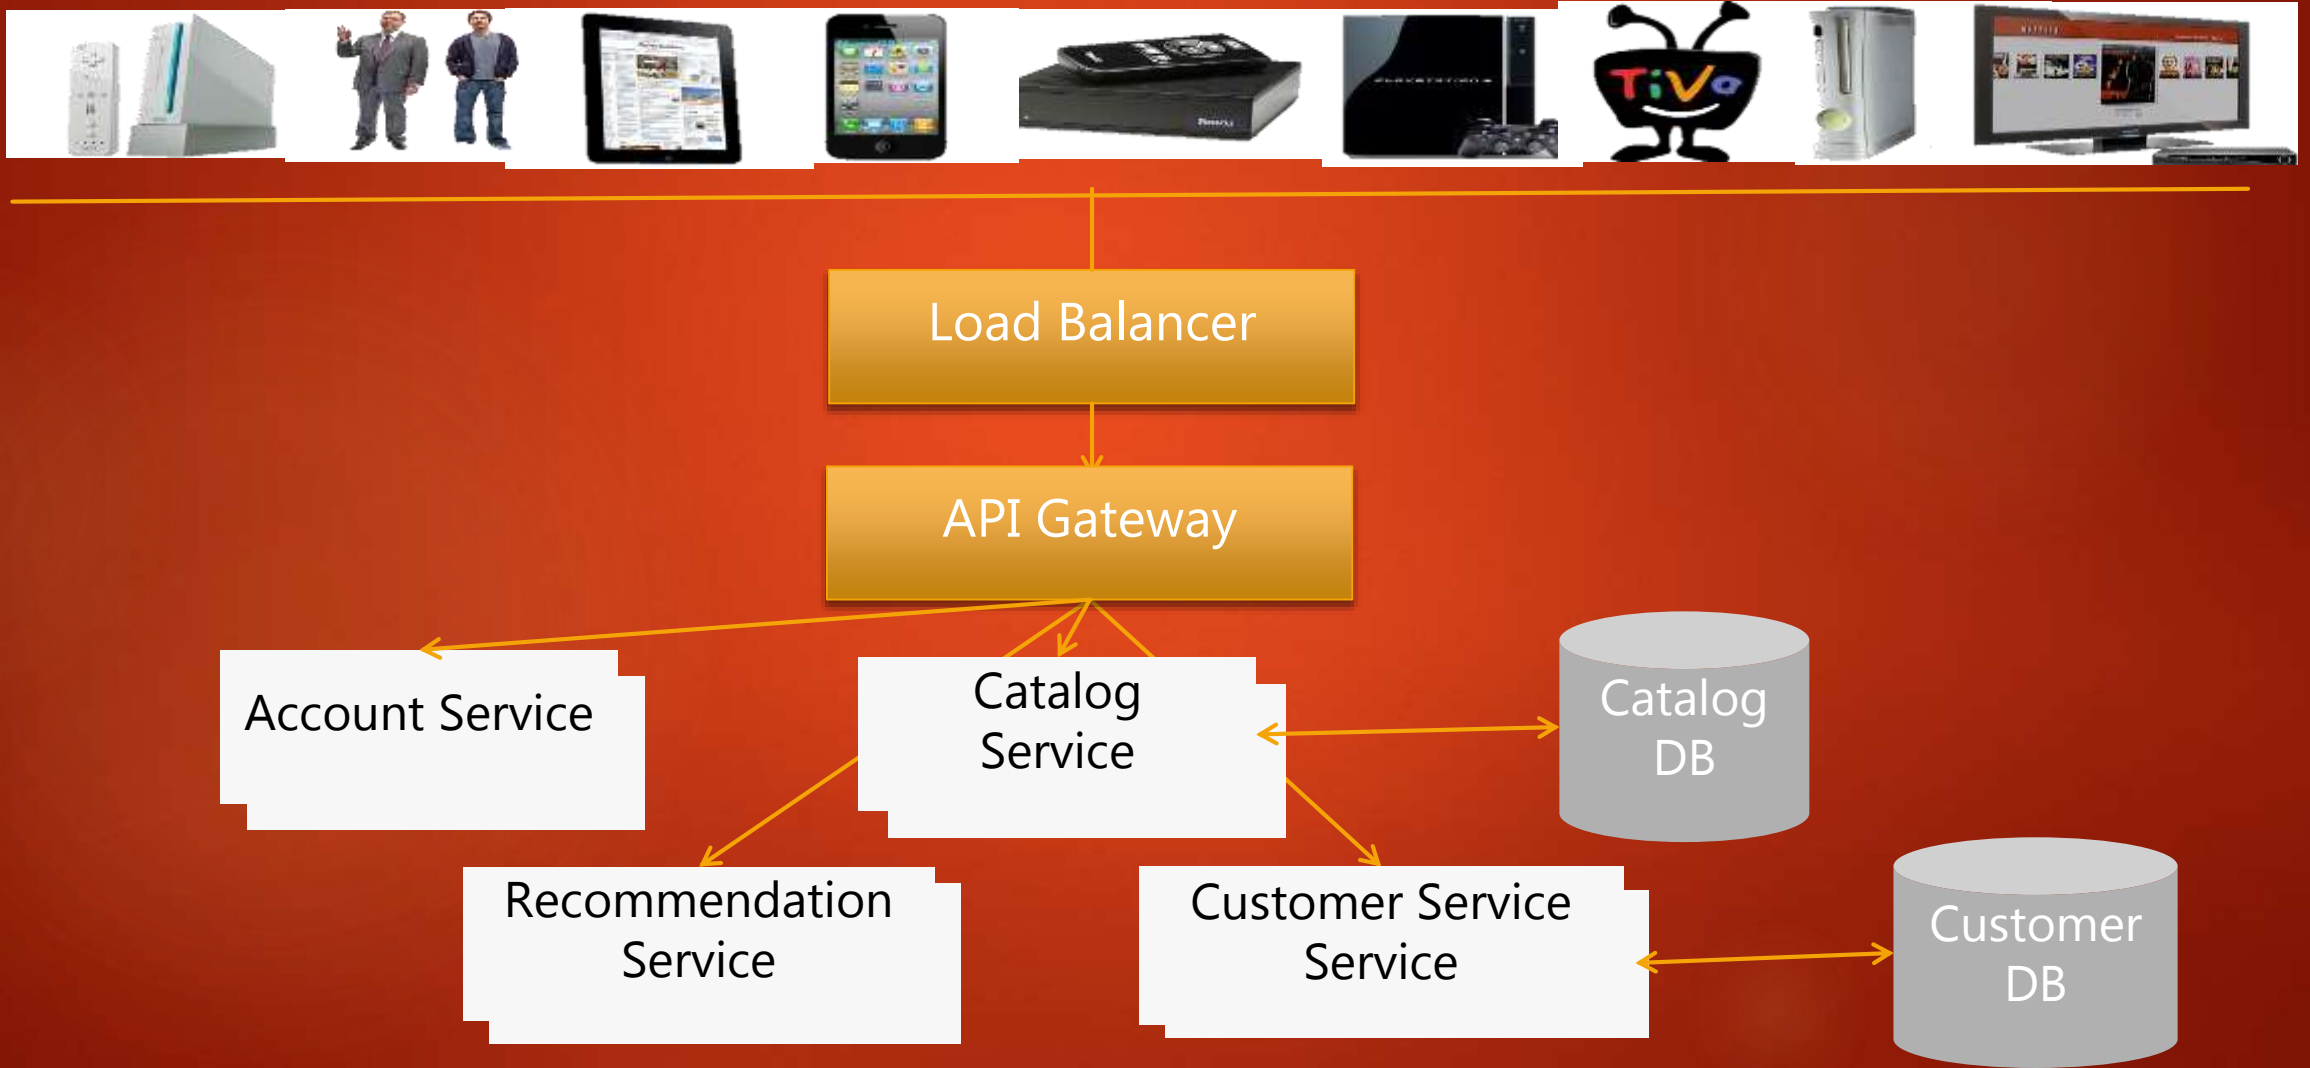
FOR THE BITEA



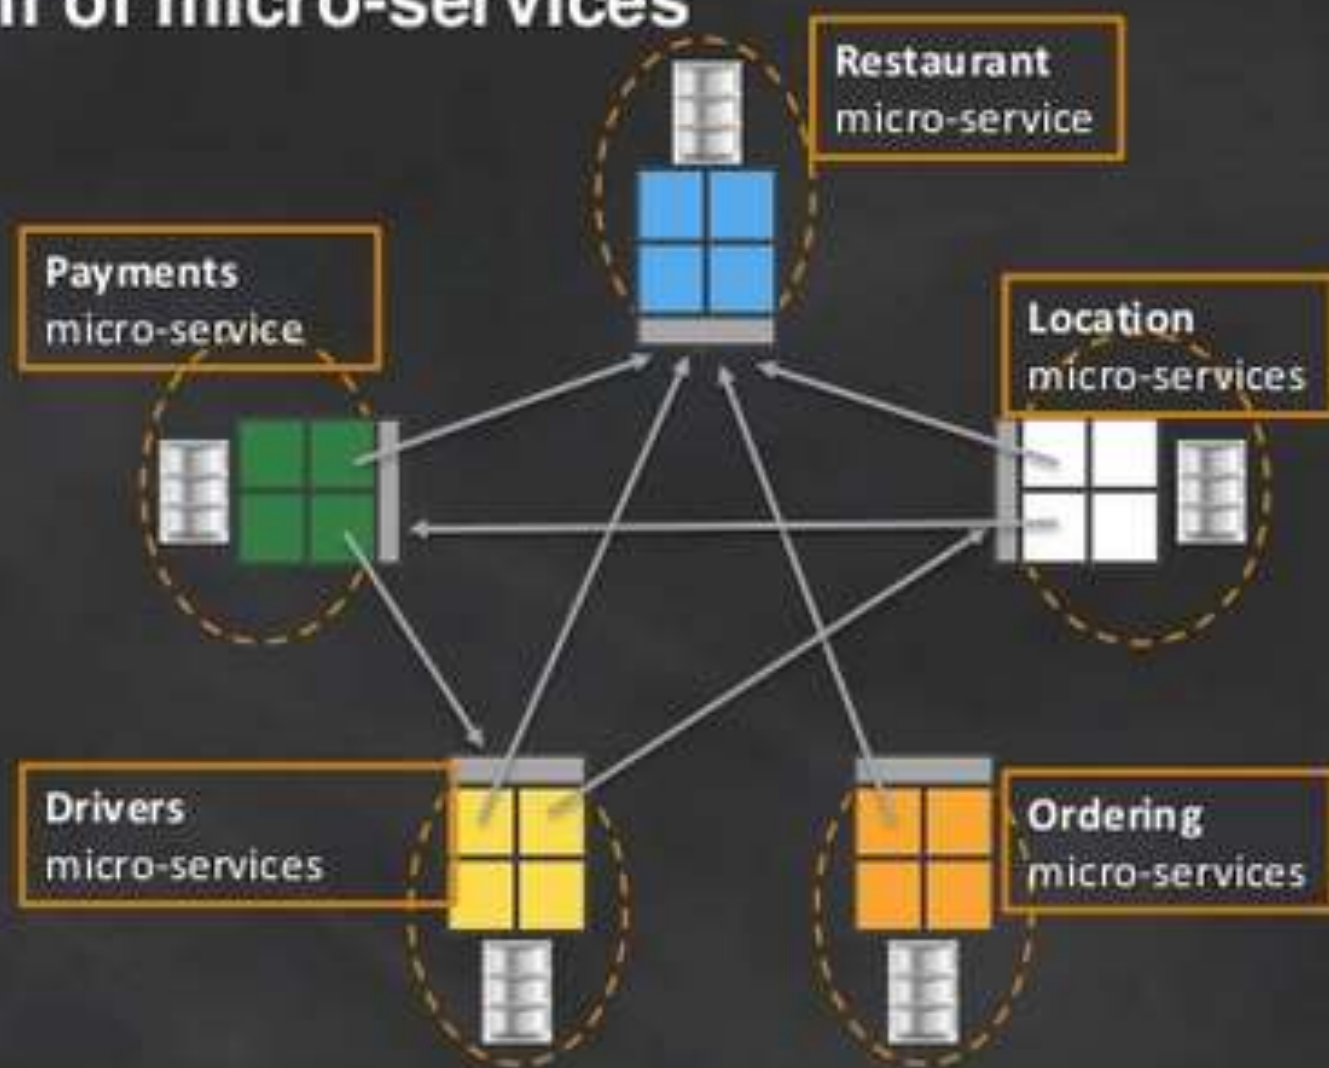


MicroServices - separate single purpose services

Microservices Architecture



Ecosystem of micro-services





what are microservices

- Micro service architecture (MSA) is an approach to building software systems that decomposes business domain models into smaller, consistent, bounded-contexts implemented by services.
- Typically implemented and operated by small teams.
- Switching from SOAP to REST doesn't make a micro services architecture.
- Micro services are not a technology-only discussion.

"service-oriented architecture

composed of
loosely coupled
elements
that have
bounded contexts"

Services communicate with
each other over the
network

*Adrian Cockcroft (former Cloud Architect at Netflix,
now Technology Fellow at Battery Ventures)*

"service-oriented
architecture
composed of
loosely coupled
elements
that have
bounded contexts"

You can update the services independently; updating one service doesn't require changing any other services.

*Adrian Cockcroft (former Cloud Architect at Netflix,
now Technology Fellow at Battery Ventures)*

"service-oriented
architecture
composed of
loosely coupled
elements
that have
bounded contexts"

*Adrian Cockcroft (former Cloud Architect at Netflix,
now Technology Fellow at Battery Ventures)*

Self-contained; you can
update the code without
knowing anything about the
internals of other
microservices

Design Principles for Monoliths:

- DDD
- SoC using MVC
- High cohesion, low coupling
- DRY
- CoC
- YAGNI

Design Patterns for Micro services:

- Aggregator Pattern
- Proxy Pattern
- Chained Pattern
- Branch Pattern
- Shared Resources
- Async Messaging and etc ..

Principle 1: Microservices only rely on each other's public API
(Hide Your Data)



Principle 2: Use the right tool for the job
(Embrace polyglot programming frameworks)



Principle 2: Use the right tool for the job
(Embrace polyglot programming frameworks)



Disadvantages of Monolith:

- Difficult to deploy and maintain
- Obstacle to frequent deployments
- Dependency between unrelated features
- Makes it difficult to try out new technologies/framework

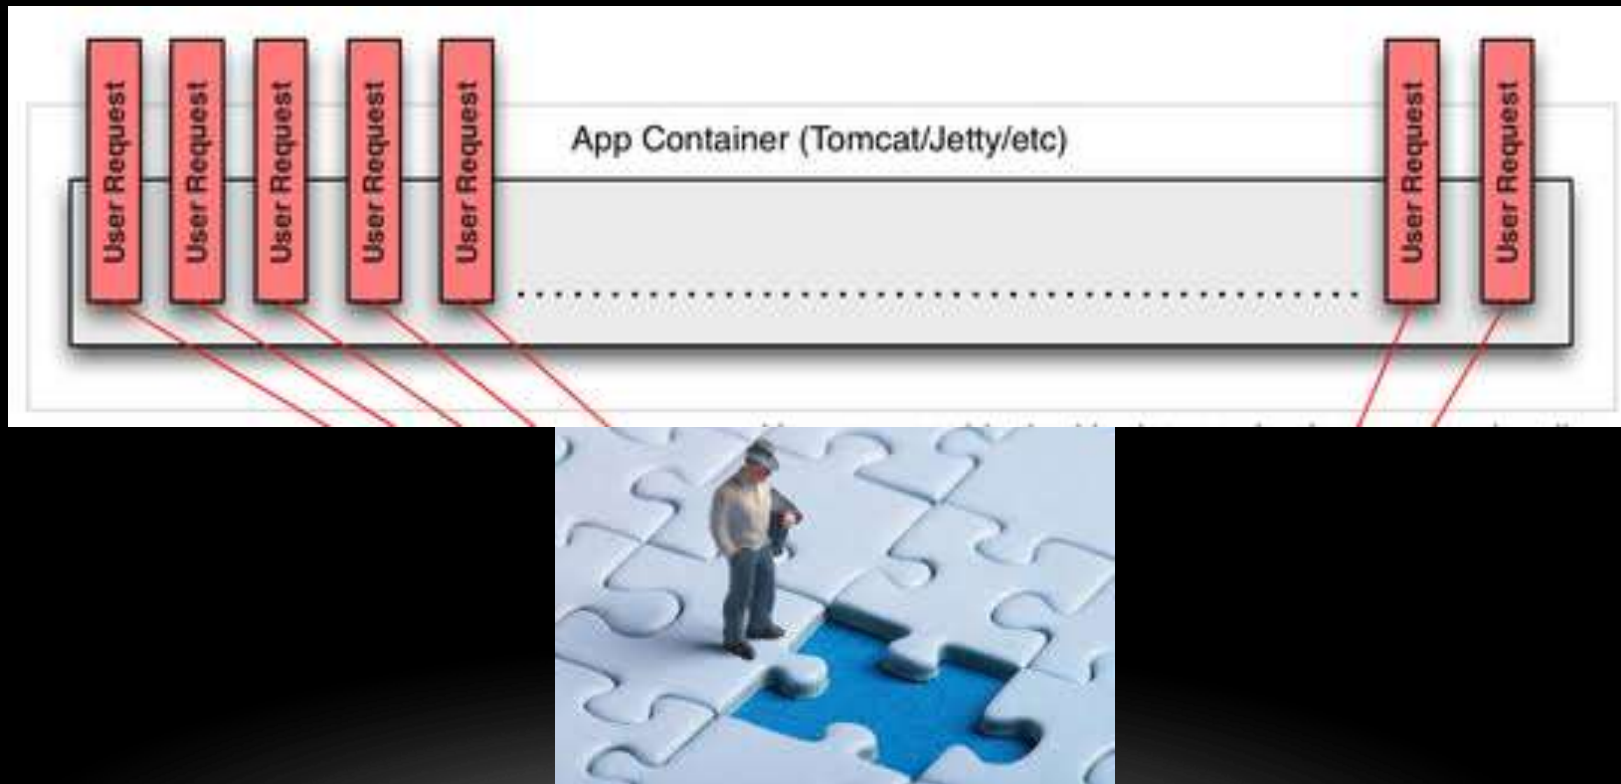
Advantages of micro services:

Easier to develop, understand, maintain

- Starts faster than a monolith, speeds up deployments
- Local change can be easily deployed, great enabler of CD
- Each service can scale on X- and Z-axis
- Improves fault isolation
- Eliminates any long-term commitment to a technology stack
- Freedom of choice of technology, tools, frameworks

Availability

- A single missing “;” brought down the Netflix website for many hours (~2008)





MONOLITHIC APPS – FAILURE & AVAILABILITY



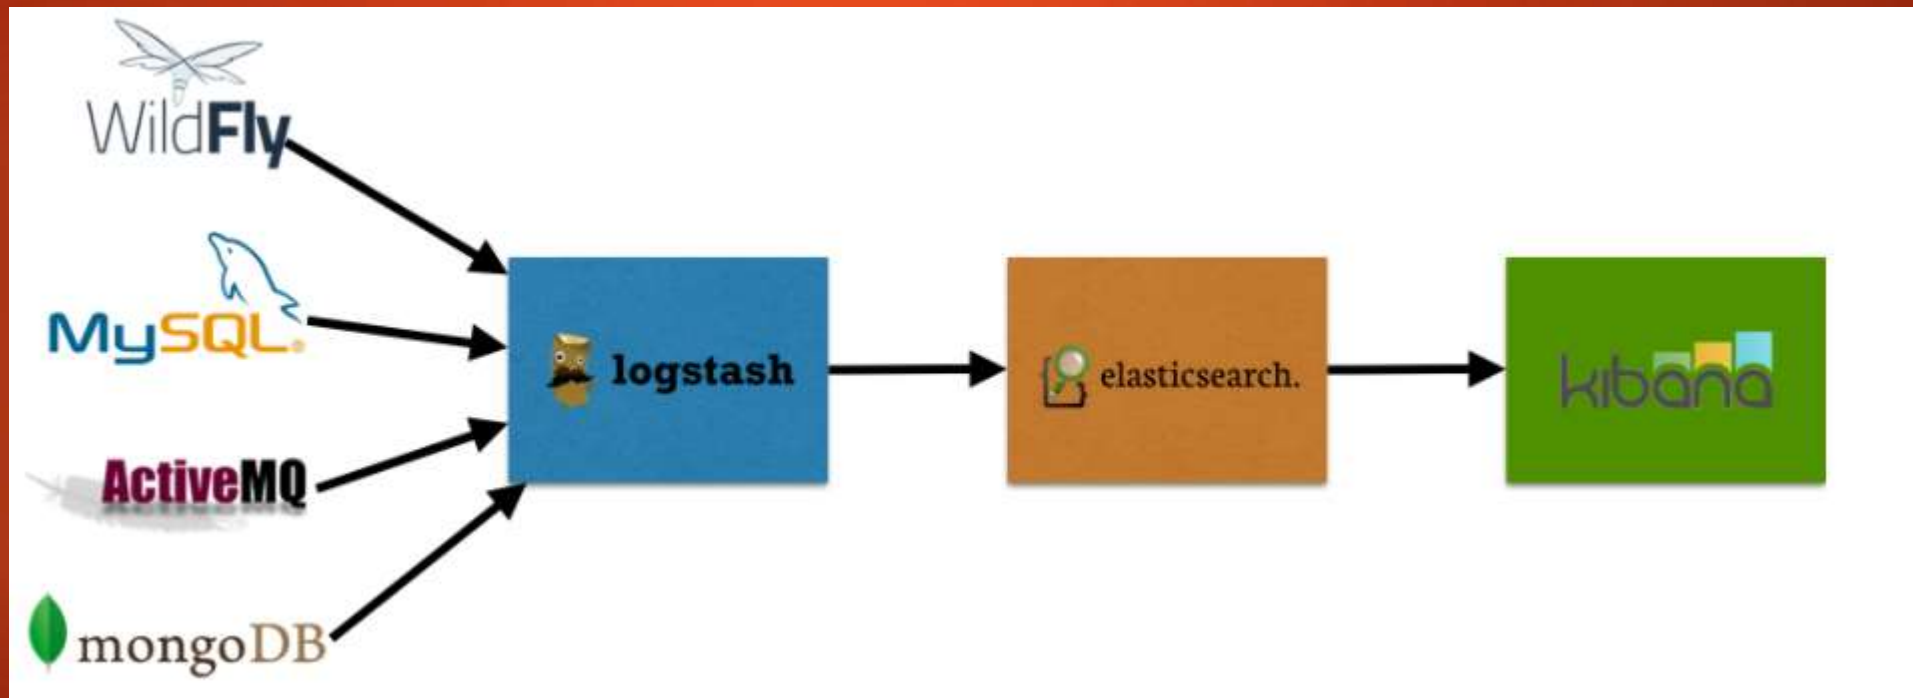
Introduction to some popular technology components, how they help solve some of the problems of developing and delivering software using a micro services architecture.

- Spring Boot (spring cloud – Eureka, ZUUL, Hystrix, Ribbon, etc)
- Docker
- Nodejs
- Kubernetes/Openshift
- NetflixOSS
- Kafka , Rabbit MQ, CloudAMQP, Kestrel etc..
- Zipkin
- ELK stack
- Prometheus with Grafana
- AppDynamics
- Akana
- ForgeRock
- API Connect etc...

Monitoring:

- One of the most frequently mentioned challenges related to the creation of micro services-based architecture is monitoring. Each micro service should be run in an environment isolated from the other micro services so it does not share resources such as databases or log files with them.
- However, the essential requirement for micro services architecture is that it is relatively easy to access the call history, including the ability to look through the request propagation between multiple micro services. Grepping the logs is not the right solution for that problem. There are some helpful tools that can be used when creating micro services with Spring Boot and Spring Cloud frameworks.
- **Zipkin**. A distributed tracing system that helps gather timing data for every request propagated between independent services. It has simple management console where we can find a visualization of the time statistics generated by subsequent services.

- **ELK.** Elasticsearch, Logstash, and Kibana — three different tools usually used together. They are used for searching, analyzing, and visualizing log data in real-time.
- Distributed, independent micro services and centralized log monitoring make for the right solution. With tools like ELK and Zipkin, microservices monitoring seems to not be a very difficult problem to solve. There are also some other tools — for example, Hystrix and Turbine Prometheus with Grafana and etc.. — that provide real-time metrics for the requests processed by micro services.



Some Tools for Service Registry/Discovery:

- Zookeeper and
- Curator
- Kubernetes
- etcd
- Consul
- OSGi
- Snoop

NoOps:

- Service replication (Kubernetes)
- Dependency resolution (Nexus)
- Failover (Circuit Breaker)
- Resiliency (Circuit Breaker)
- Service monitoring, alerts and events (ELK)

Open Source

➤ Commoditization of Technology

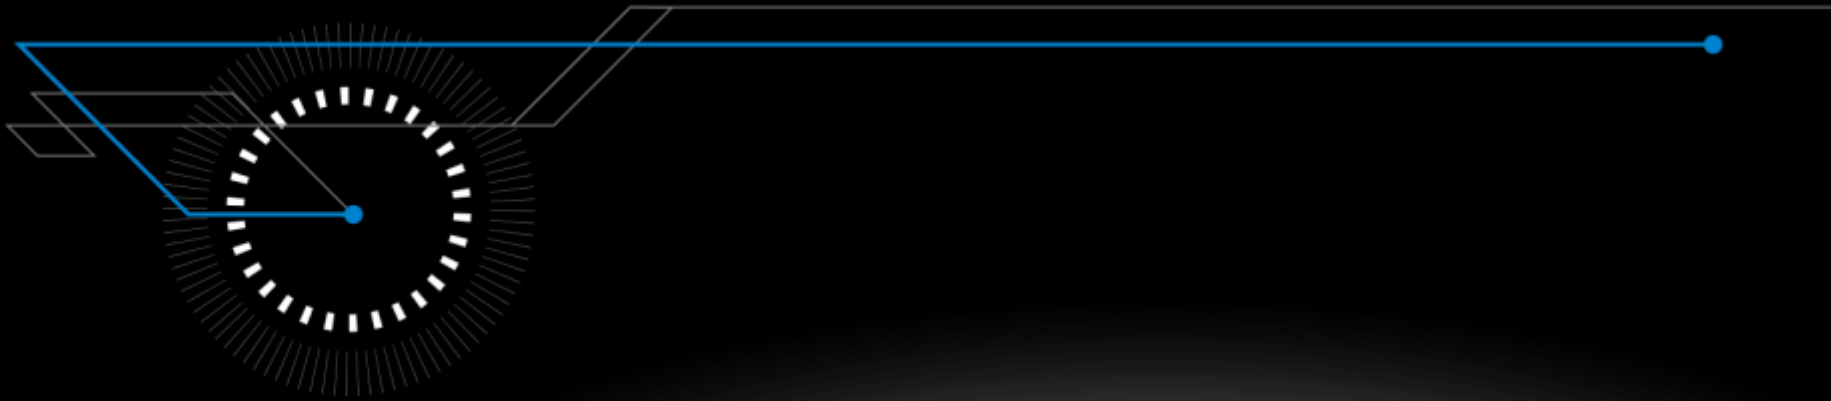
Open source is leading the charge in the technology space.

This drives communities to build things like operating systems (Linux), programming languages (Go), message queues (Apache ActiveMQ), and web servers (httpd).

As open source and open ecosystems have become the norm, we're starting to see a lot of innovation in software technology coming directly from open source communities (e.g., Apache Spark, Docker, and Kubernetes, ELK Stack, Zipkin).

Micro Services

- Challenges



CHALLENGES



Can lead to chaos if not designed right ...

CHALLENGES

Requires thinking differently about how to build, deploy and operate them.

- Design for Faults

Building distributed systems is different from building shared memory, single process, monolithic applications.

Networks are inherently unreliable.


Latent network calls can be very difficult to debug.

- Design with Dependencies in Mind

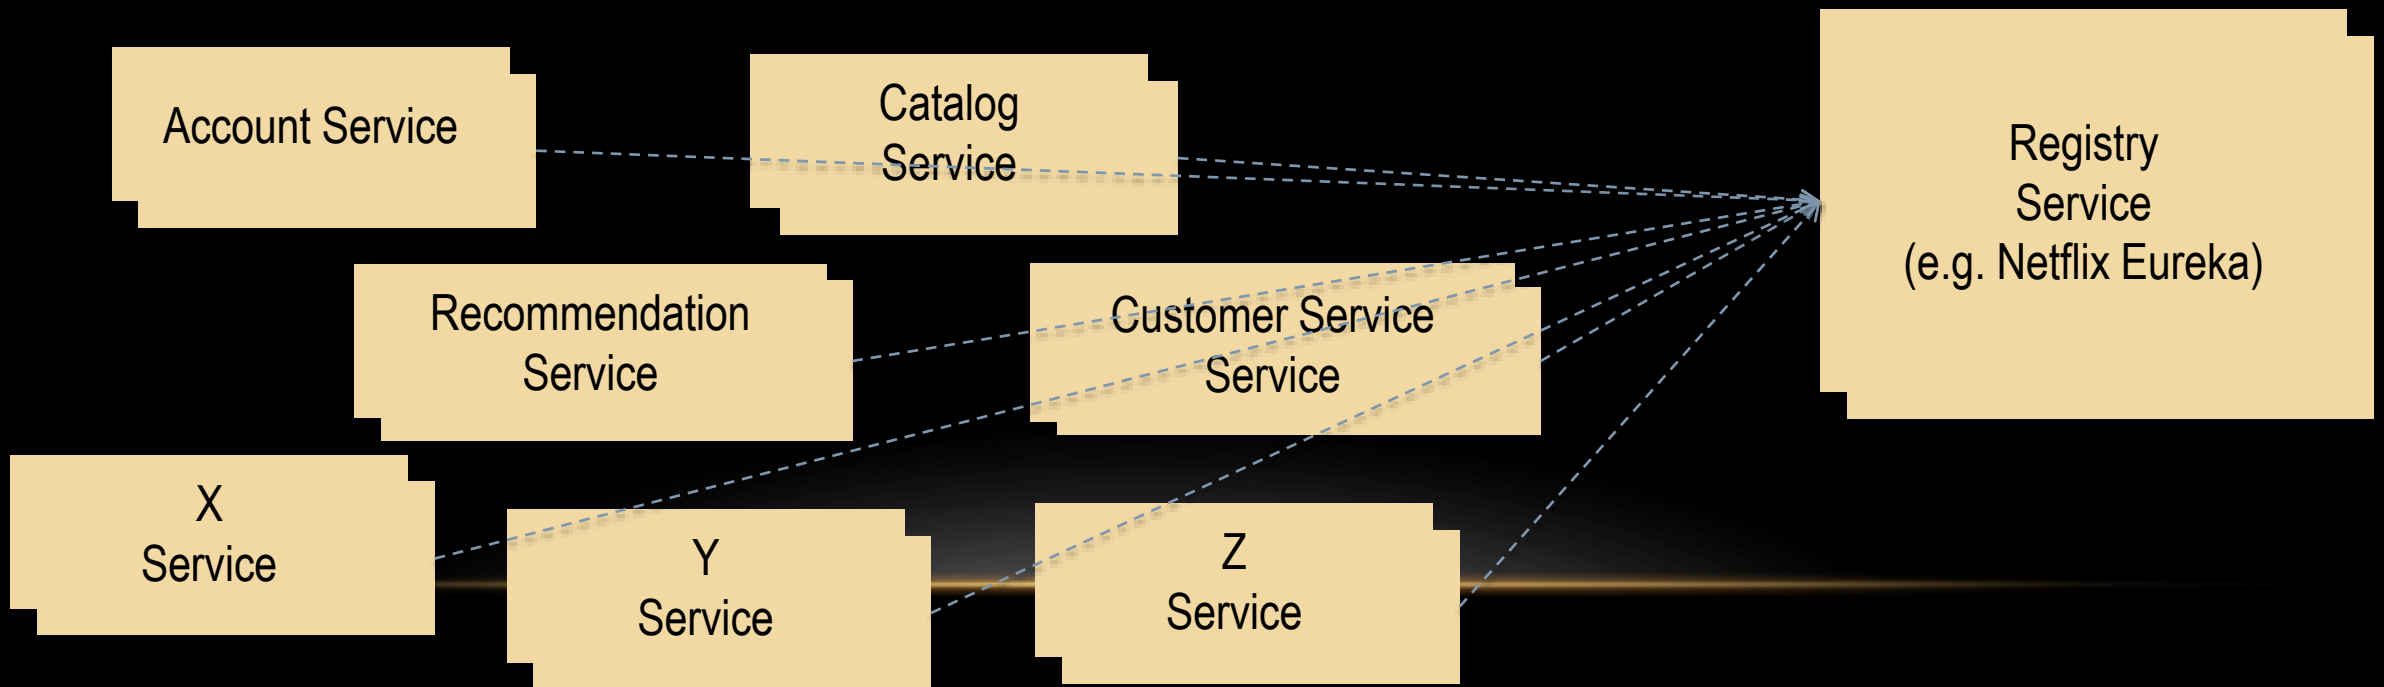
Need loose coupling in our teams, in our technology, and our governance.

SERVICE DISCOVERY

- 100s of MicroServices
 - Need a Service Metadata Registry (Discovery Service)





NETFLIX | OSS




CHATTINESS (AND FAN OUT)

~2 Billion Requests per day on Edge Service

Results in ~20 Billion Fan out requests in ~100 MicroServices



Best Practices/Tips



Best Practice -> Loadbalancers

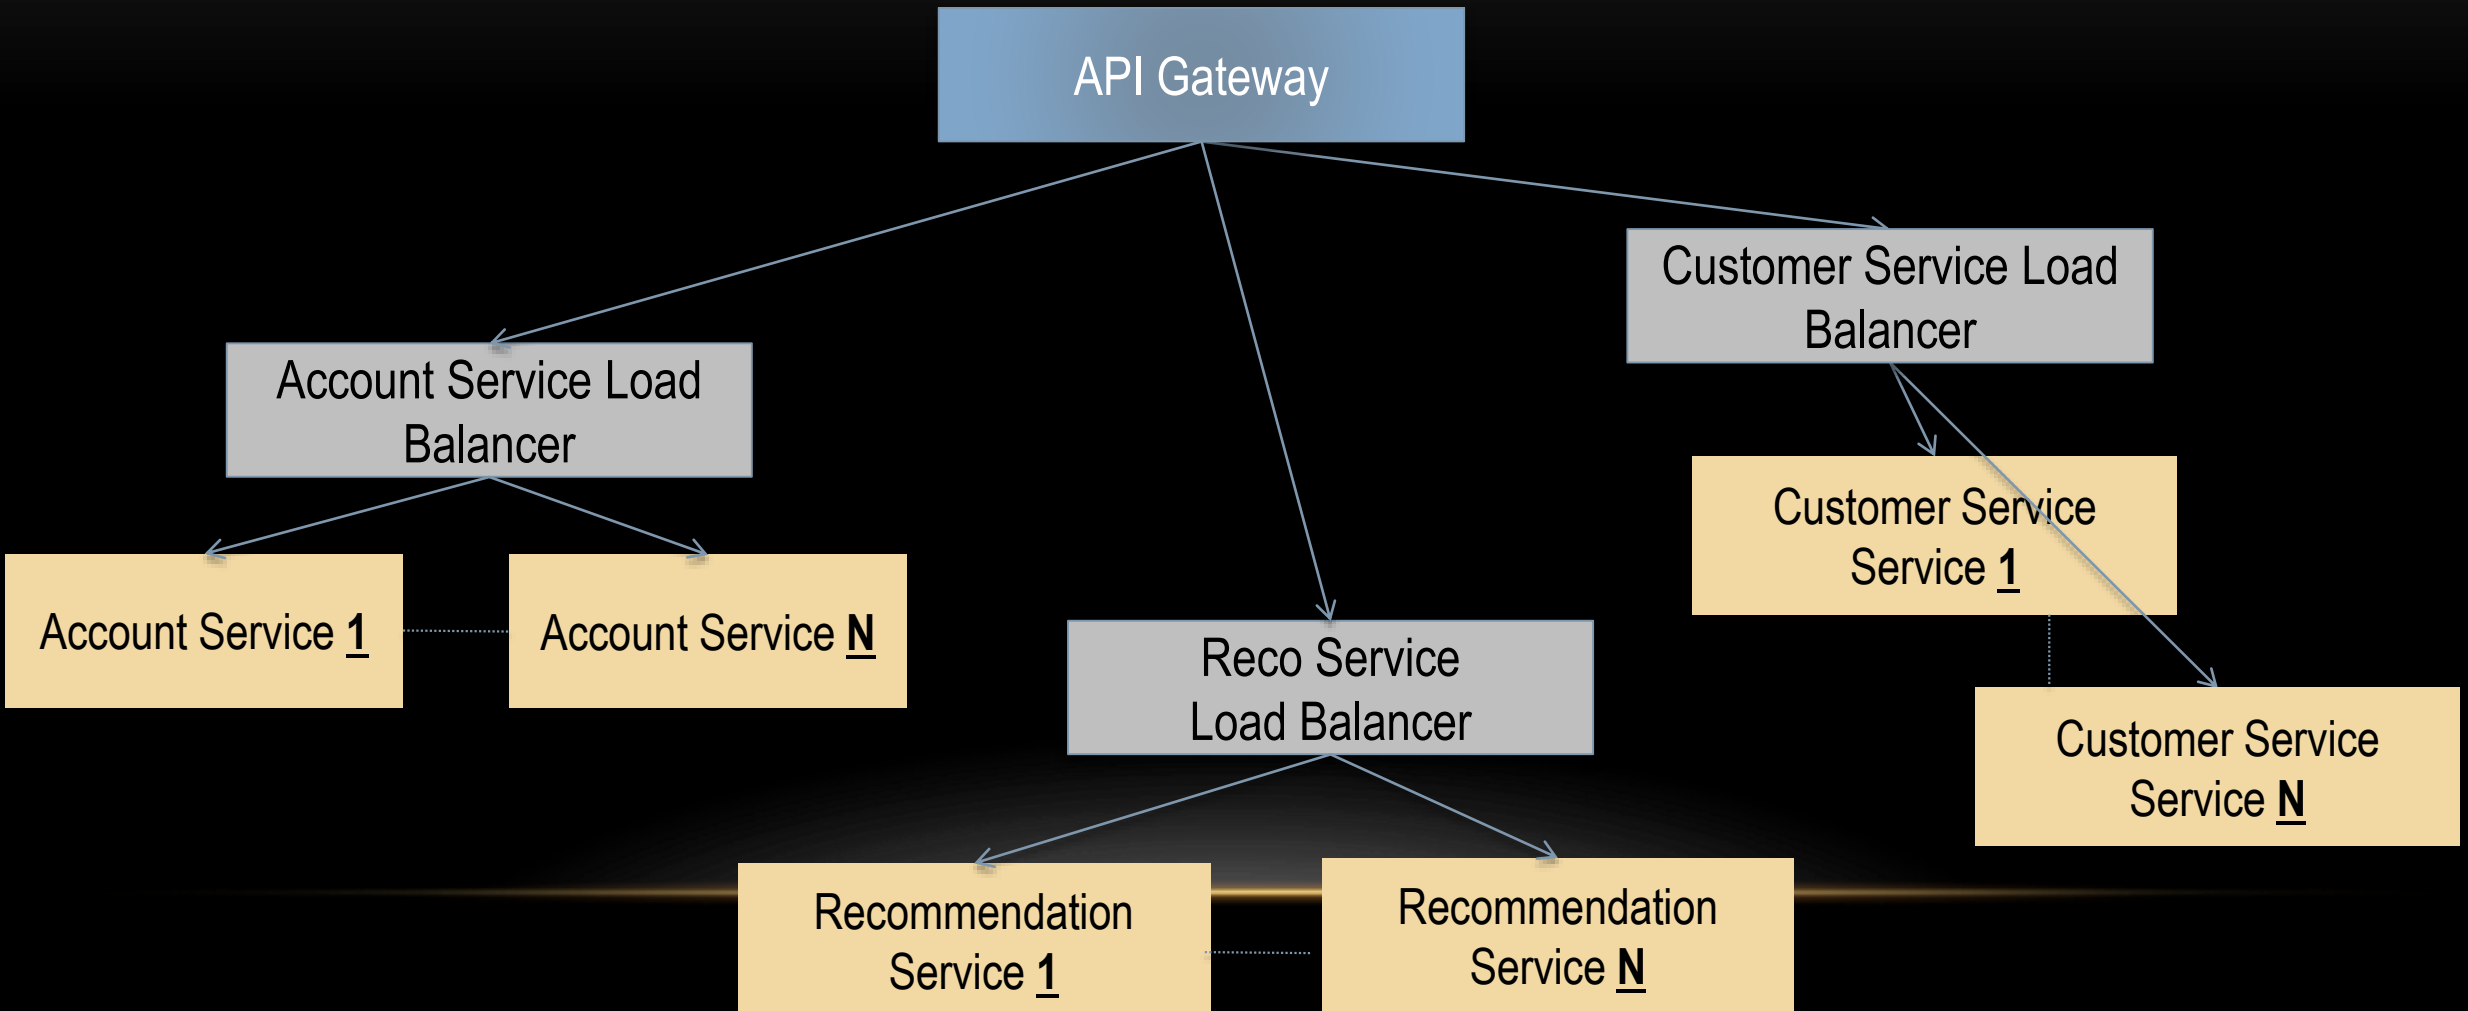
Choice

1. Central Loadbalancer?

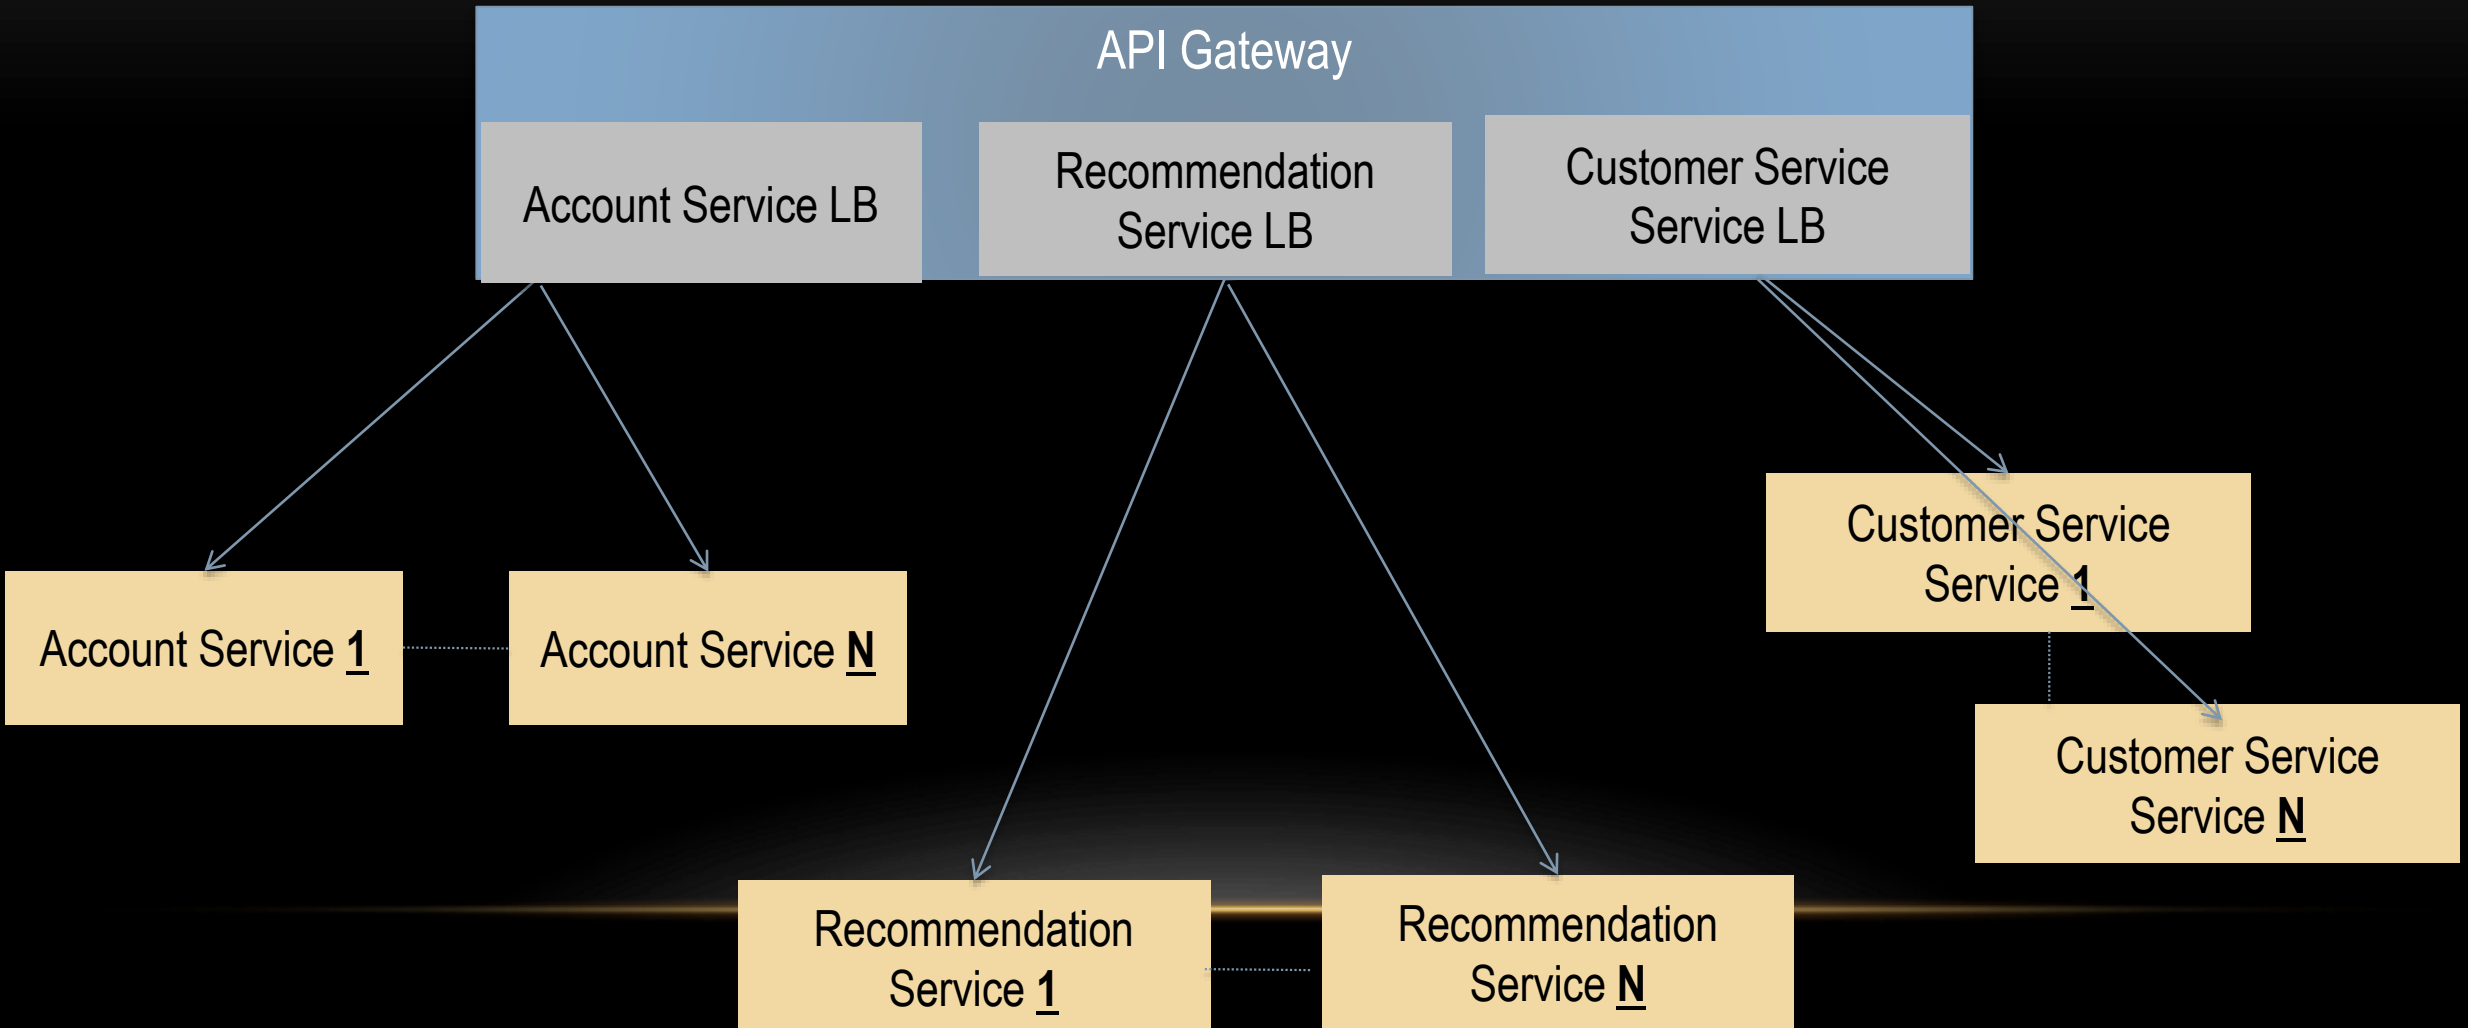
OR

2. Client based Loadbalancer?

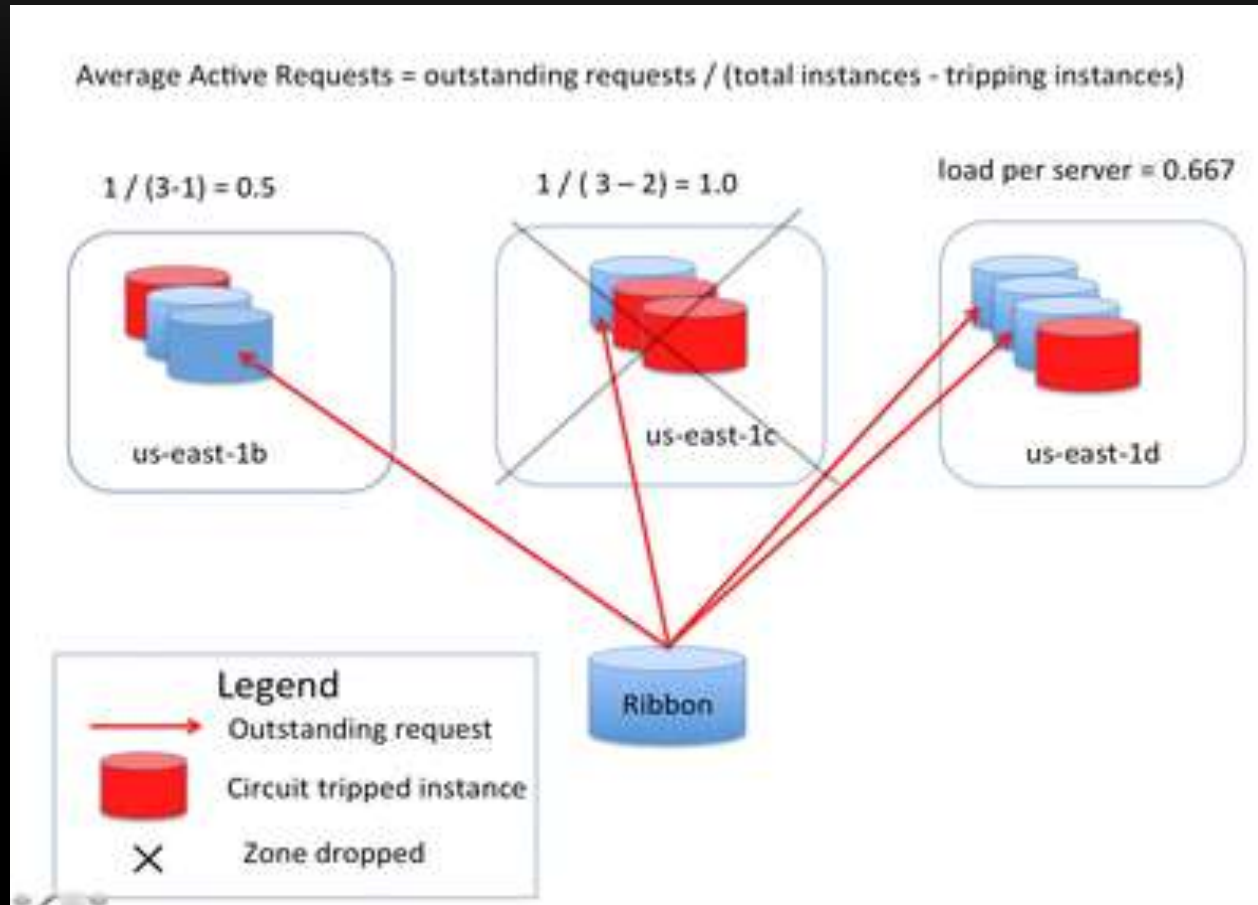
Central (Proxy) Loadbalancer



Client Loadbalancer

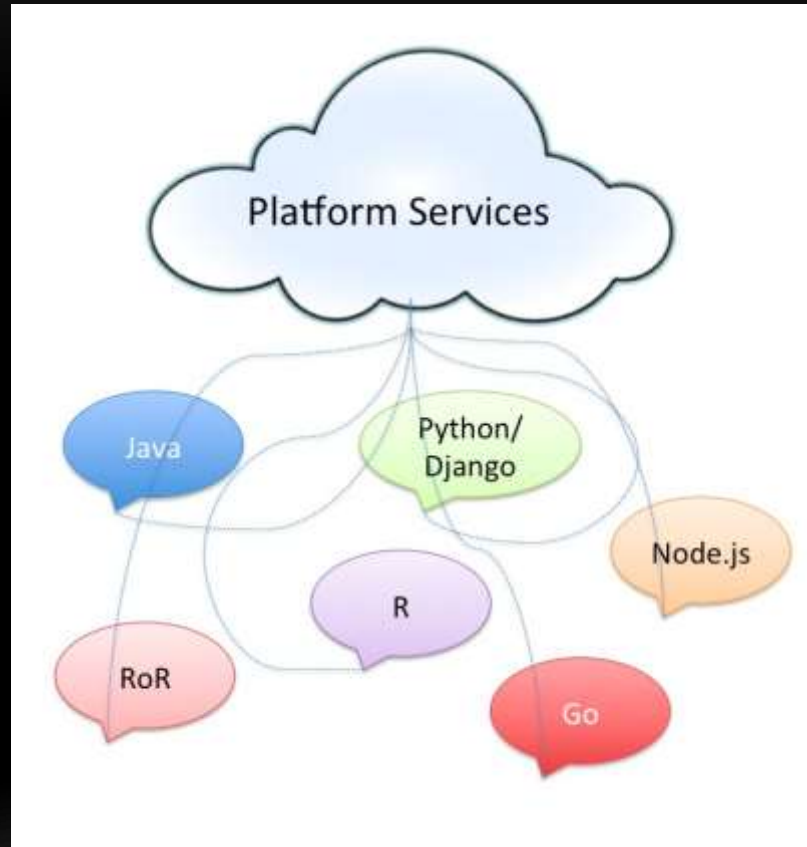


Client based Smart Loadbalancer



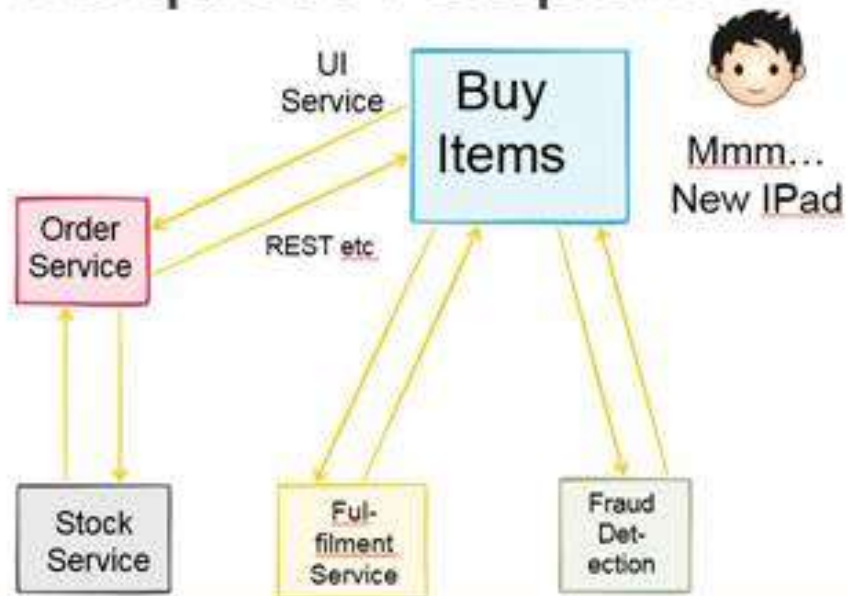
Use Ribbon (<http://github.com/netflix/ribbon>)

Homogeneity in A Polyglot Ecosystem

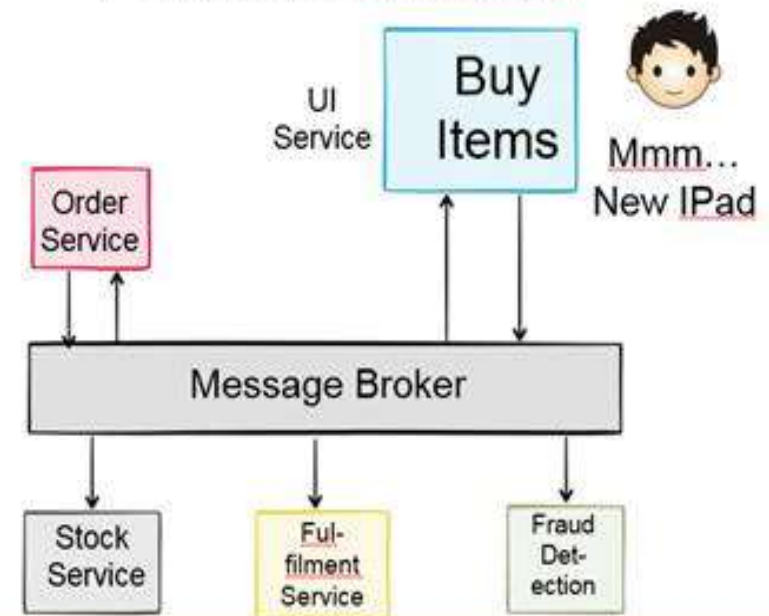


Message Queue

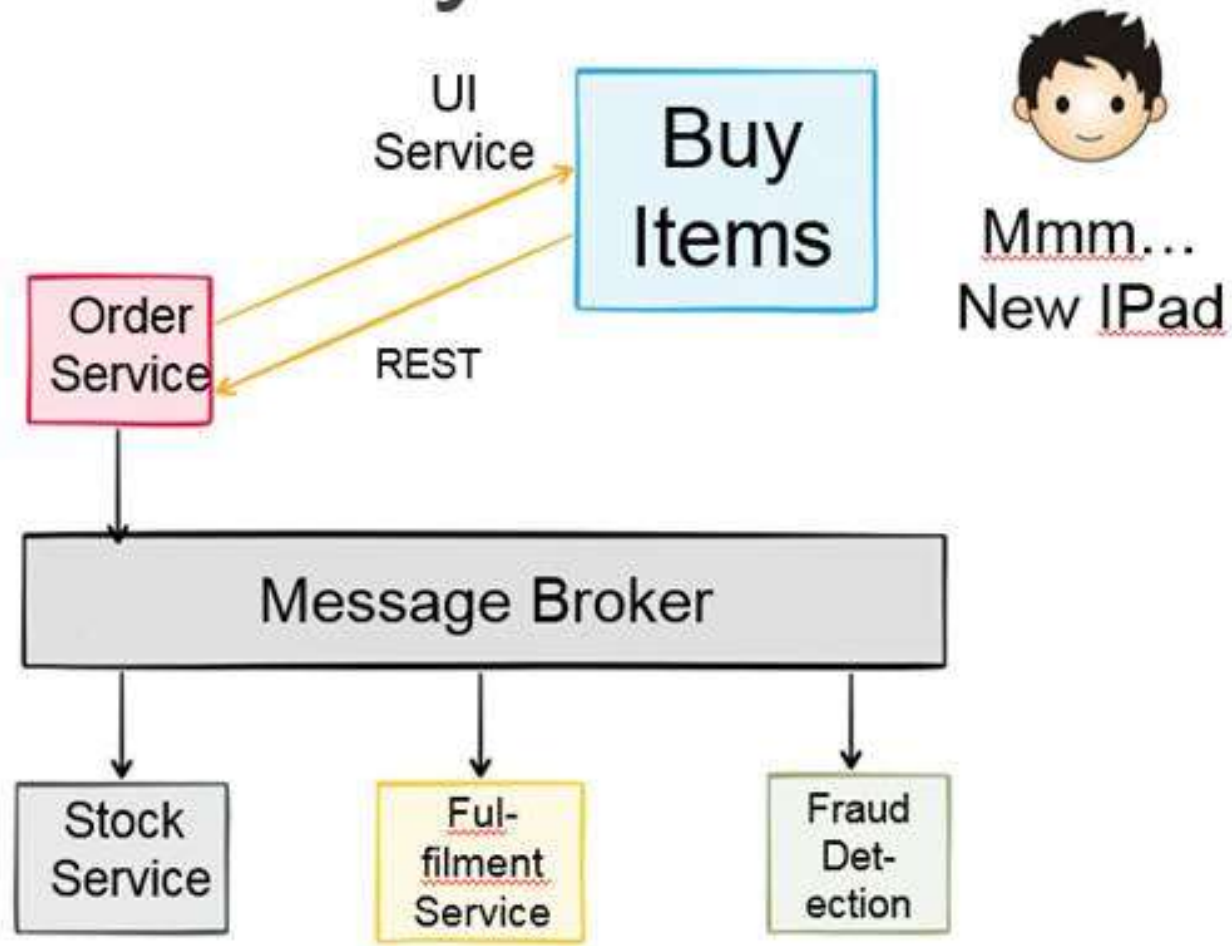
Request Response

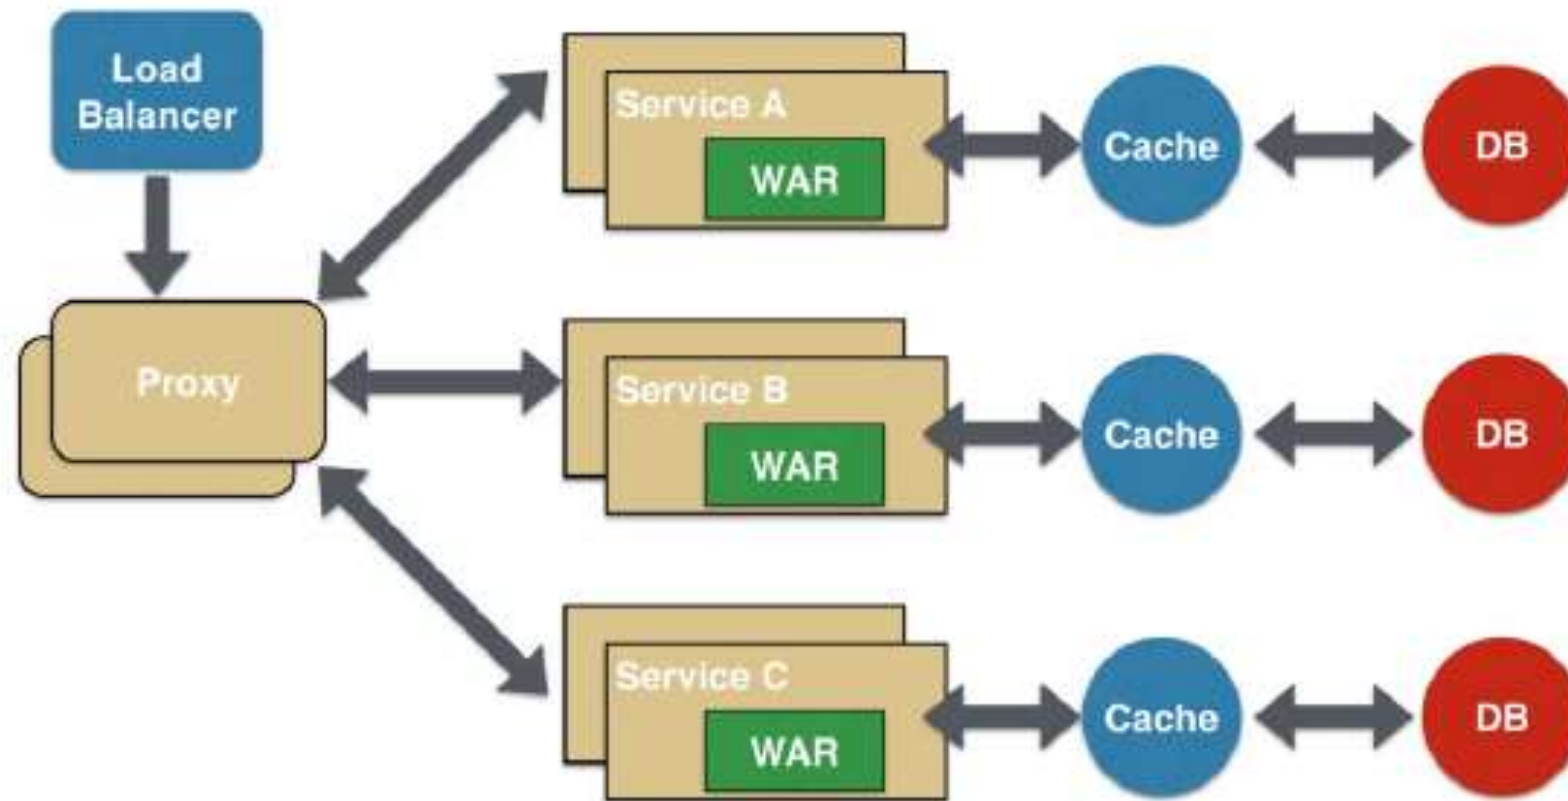


Event Driven

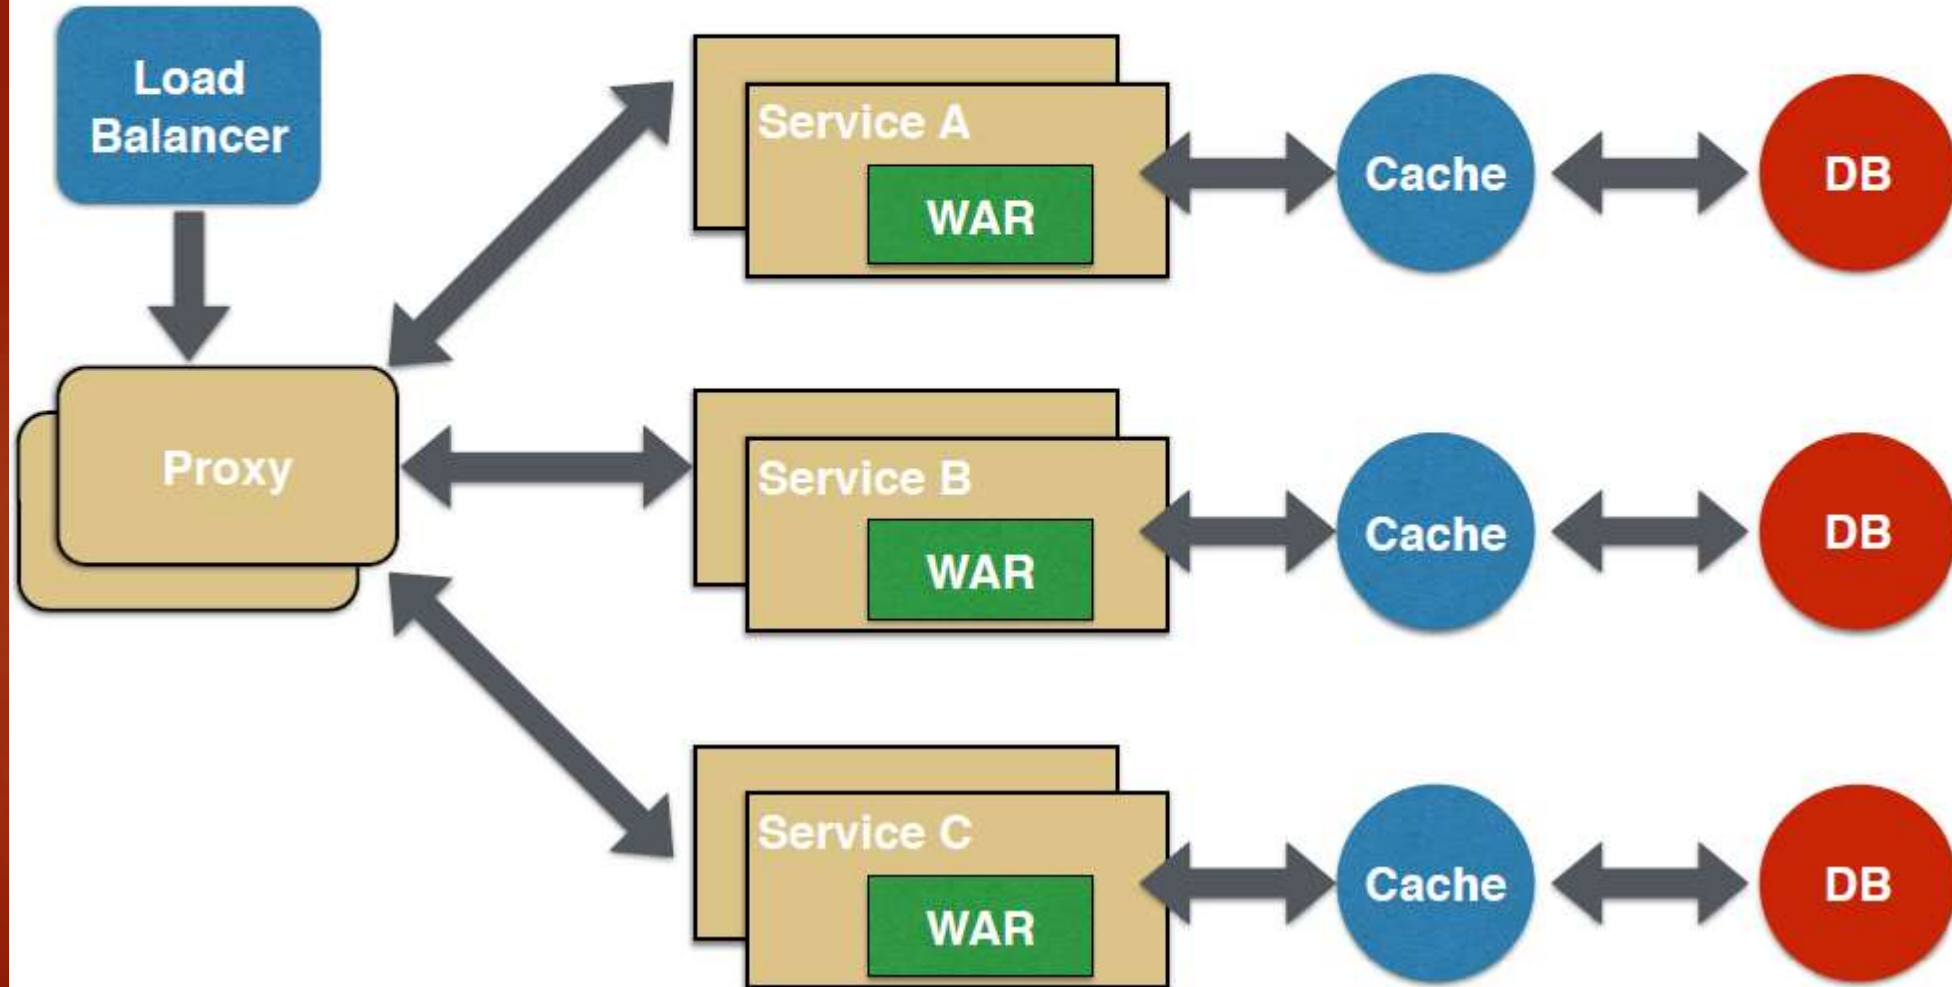


Hybrid

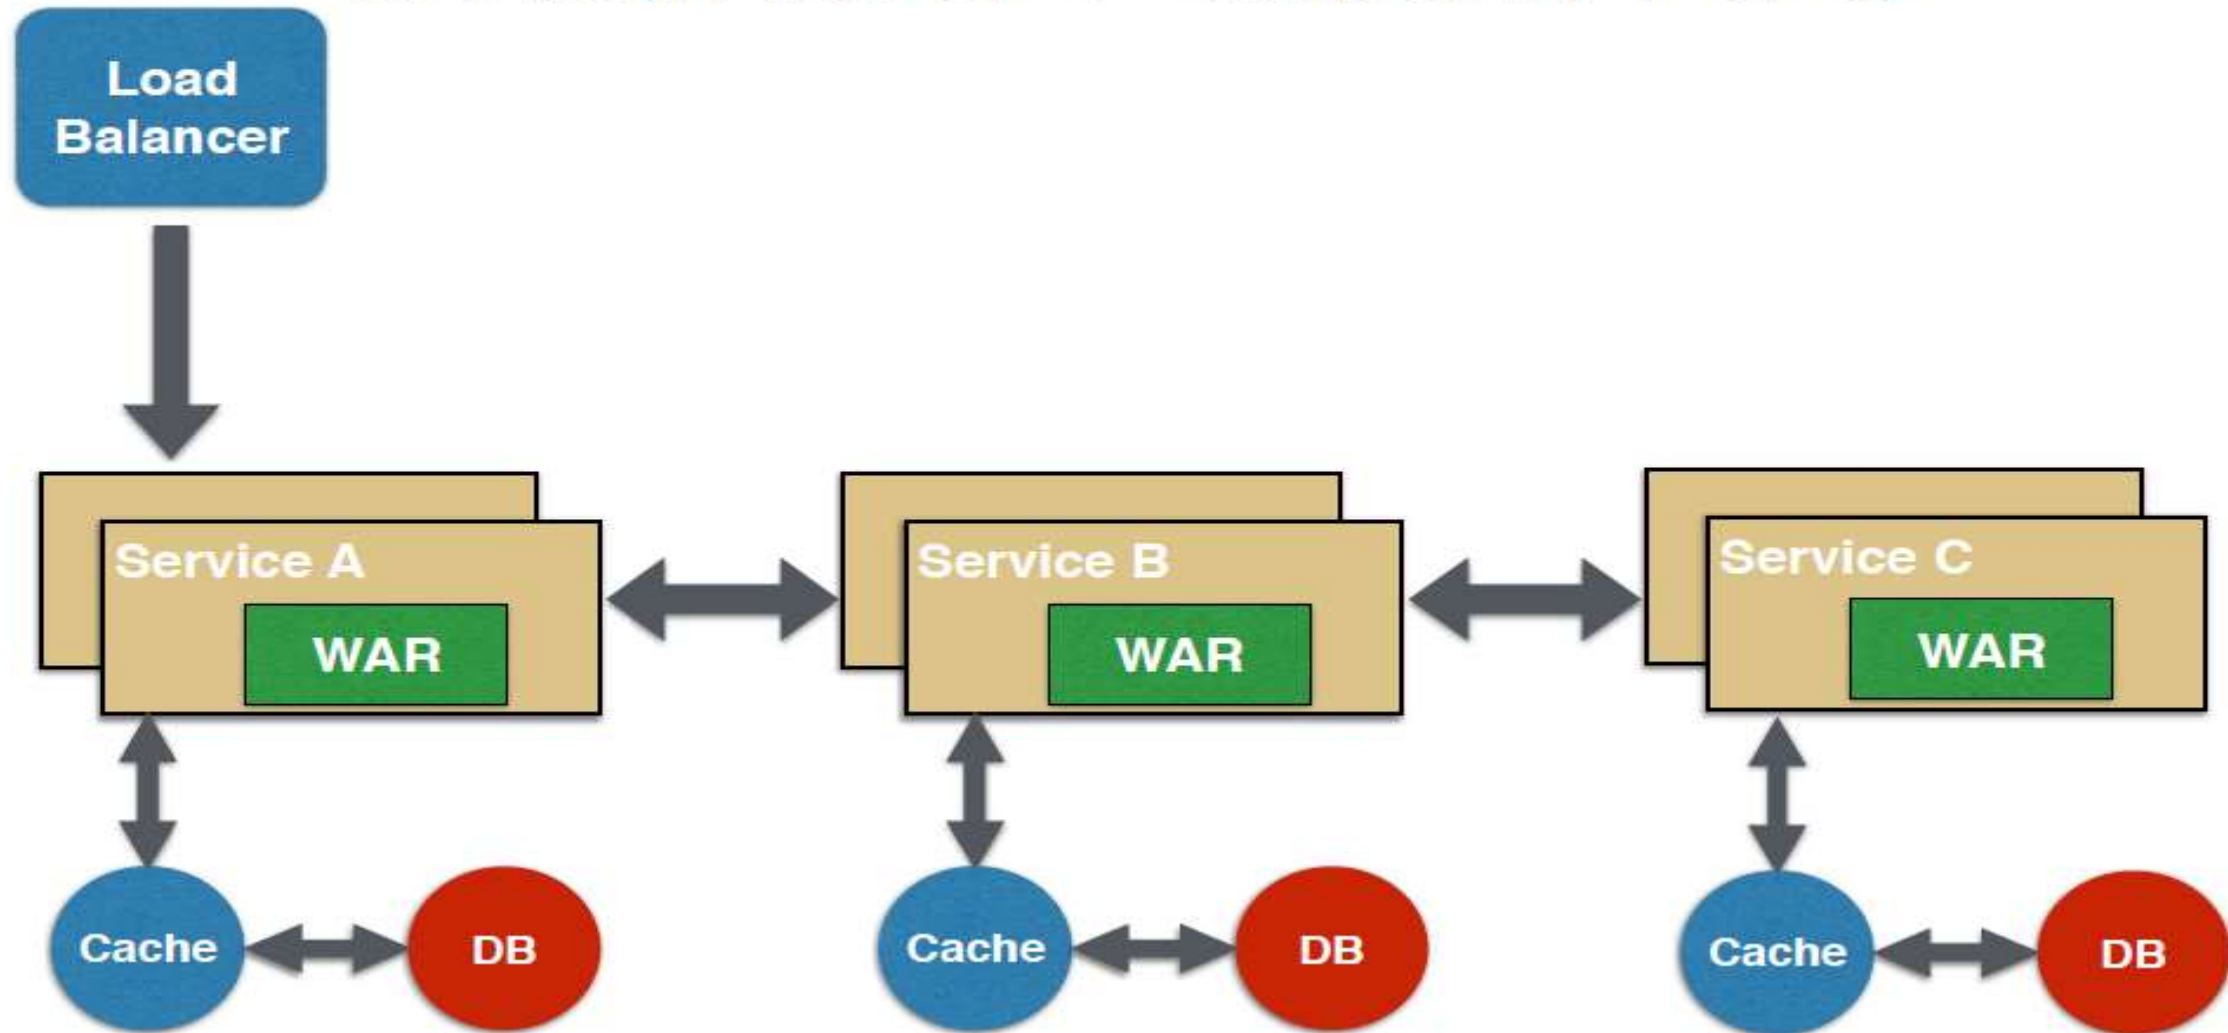




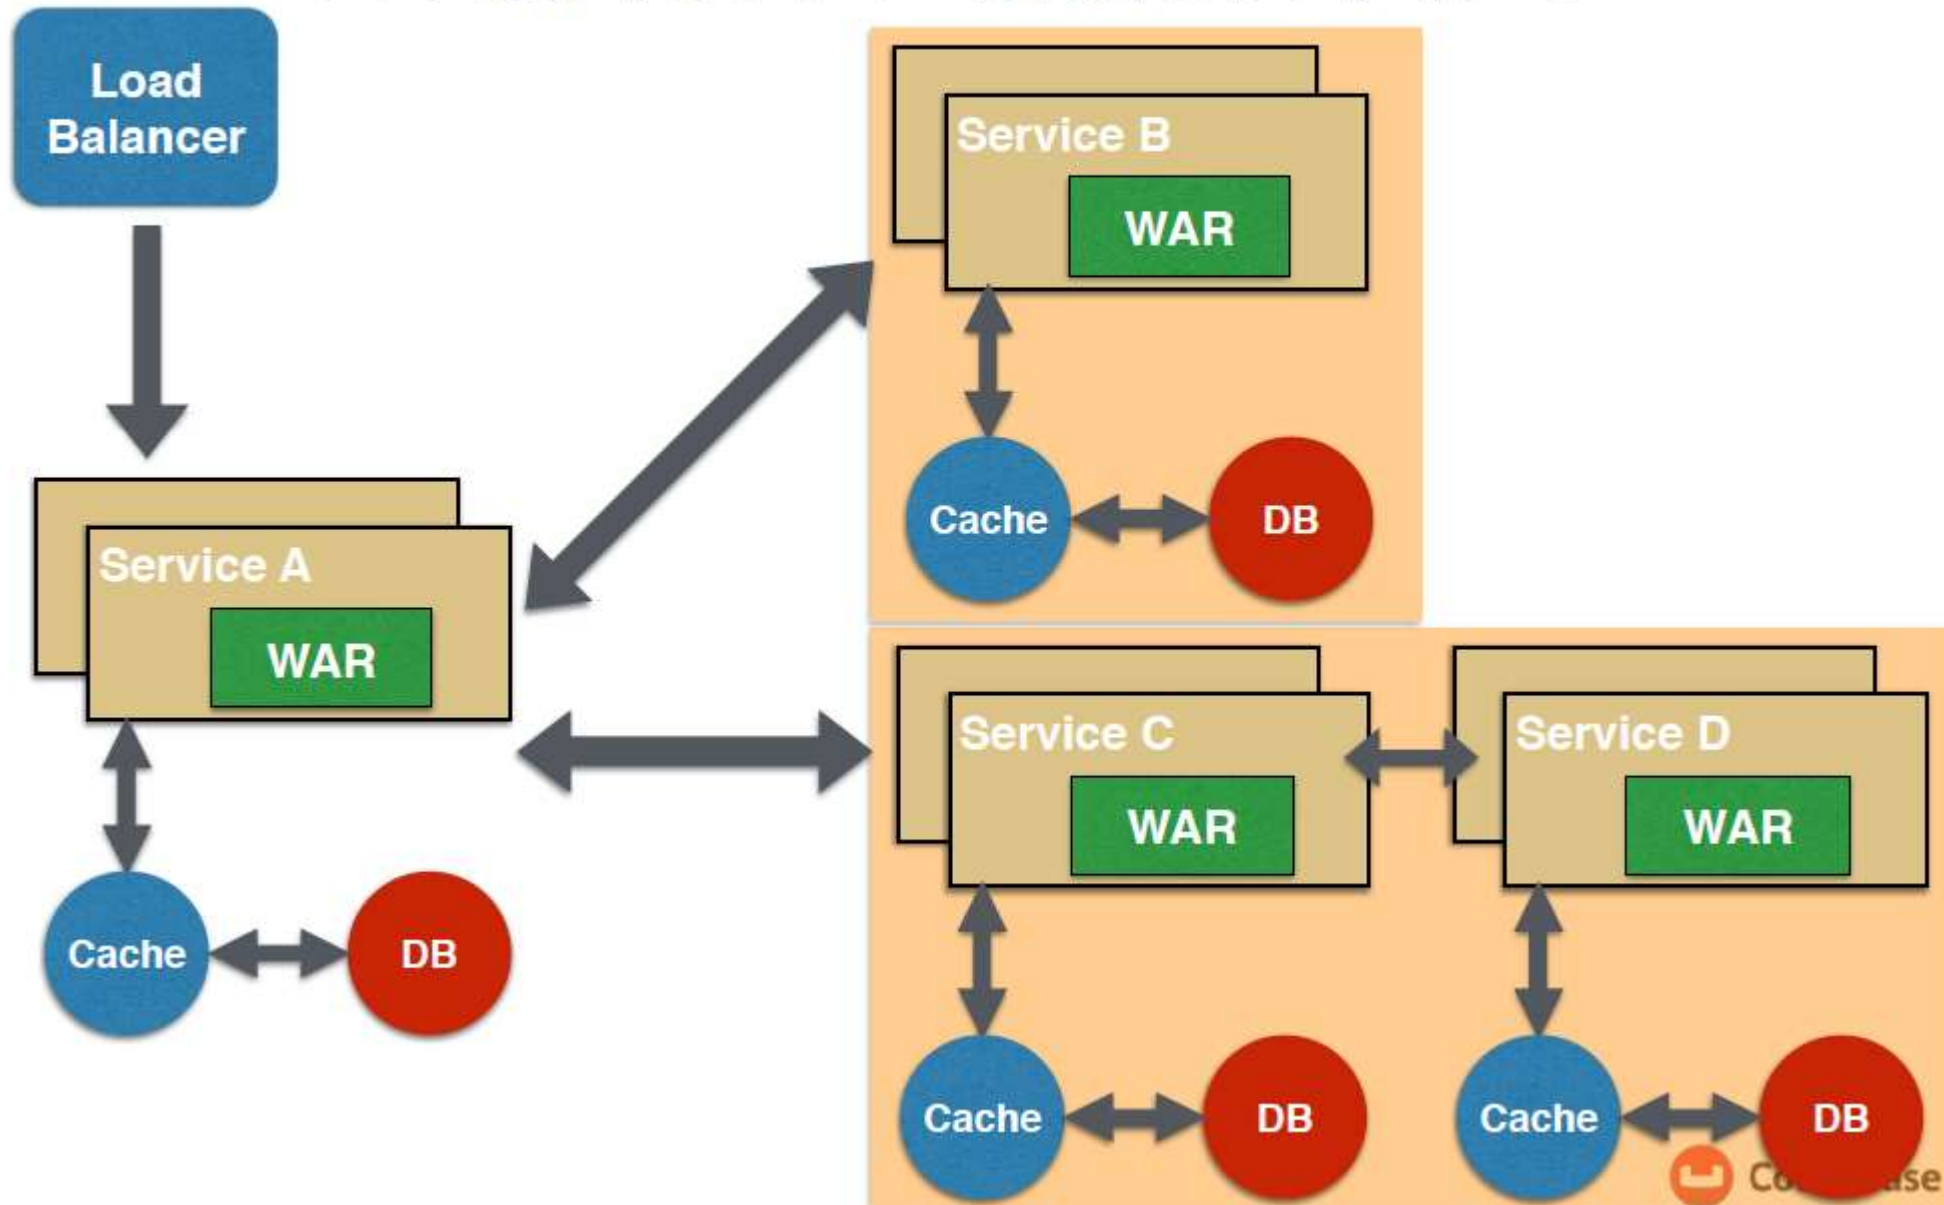
Proxy Pattern #2



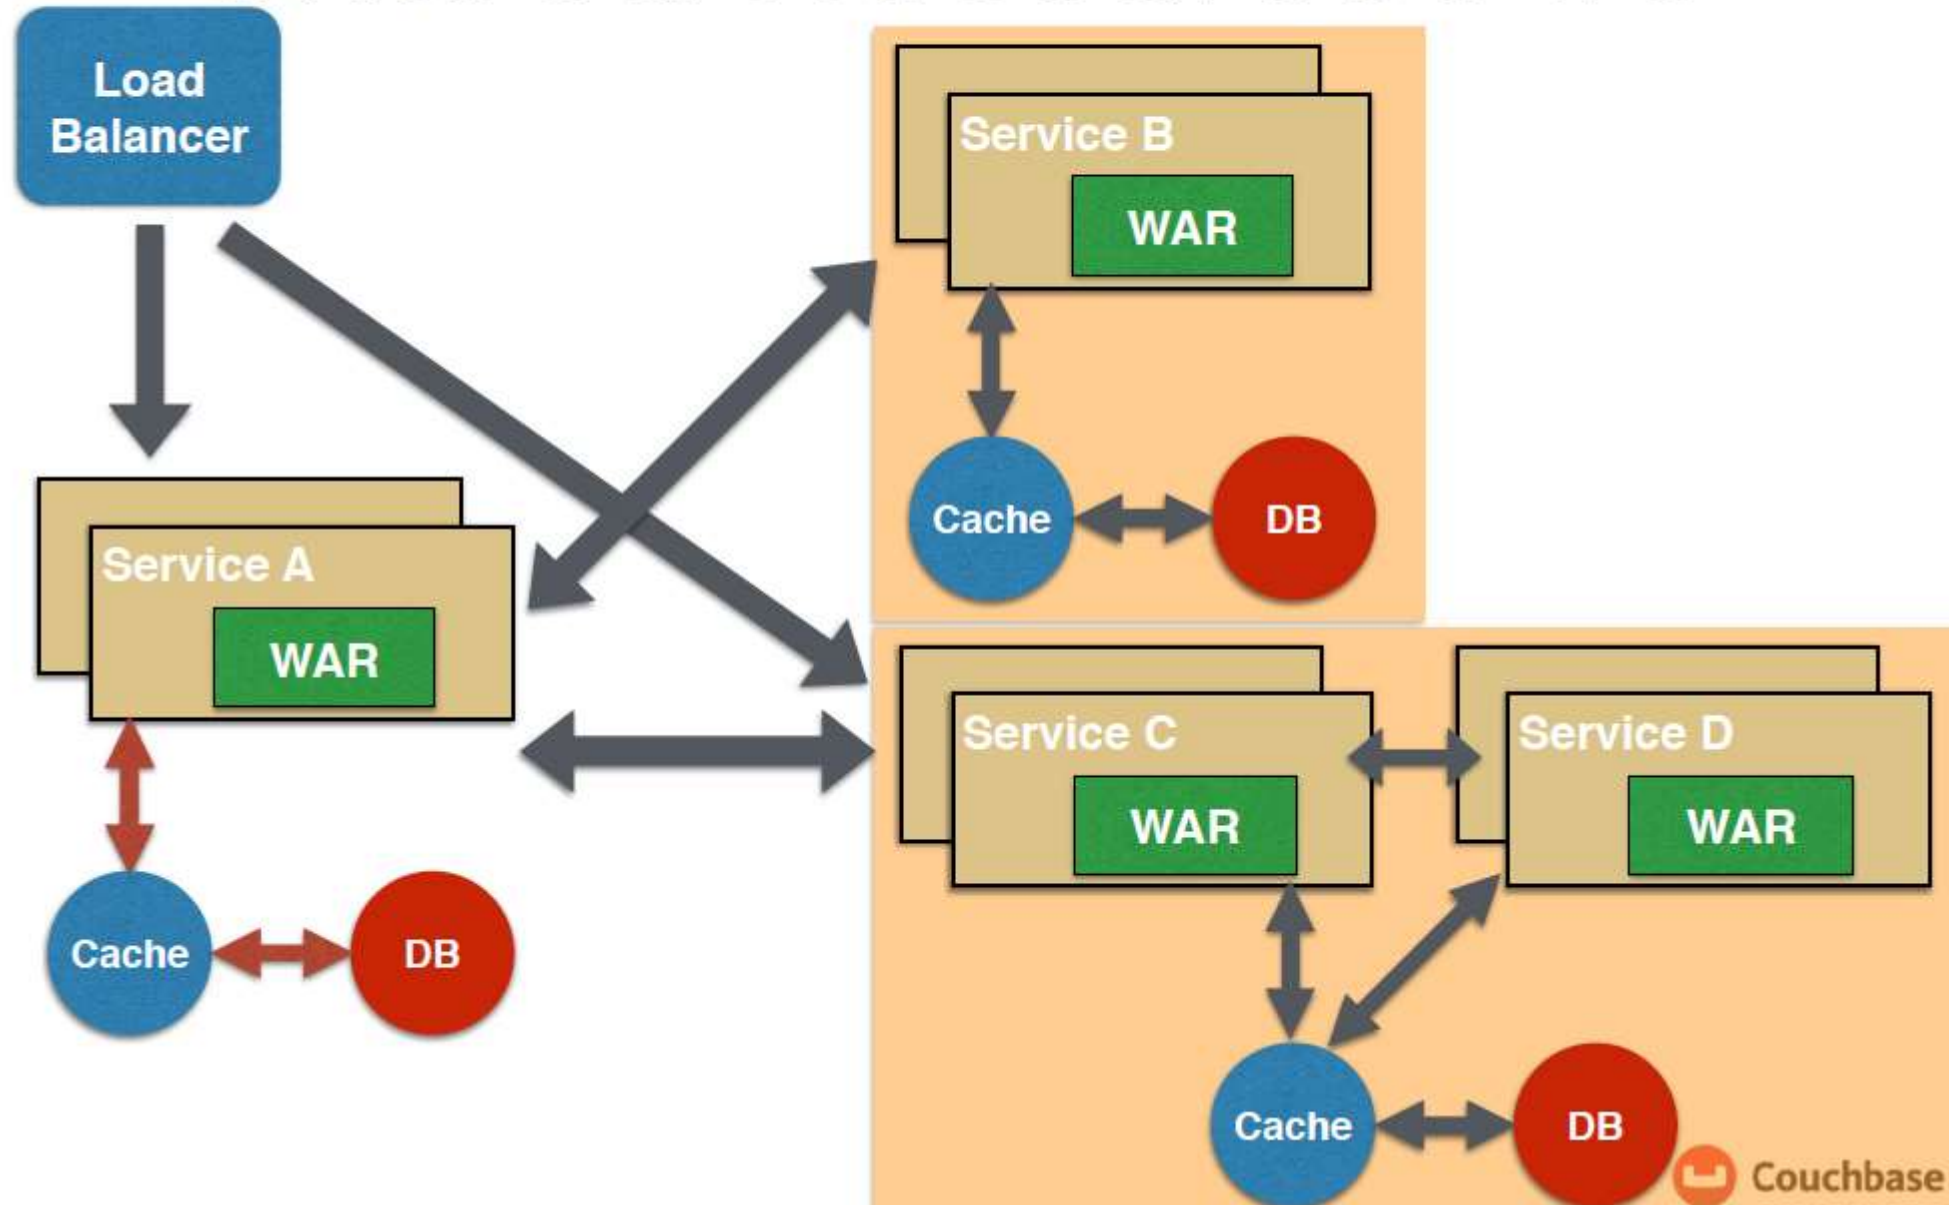
Chained Pattern #3



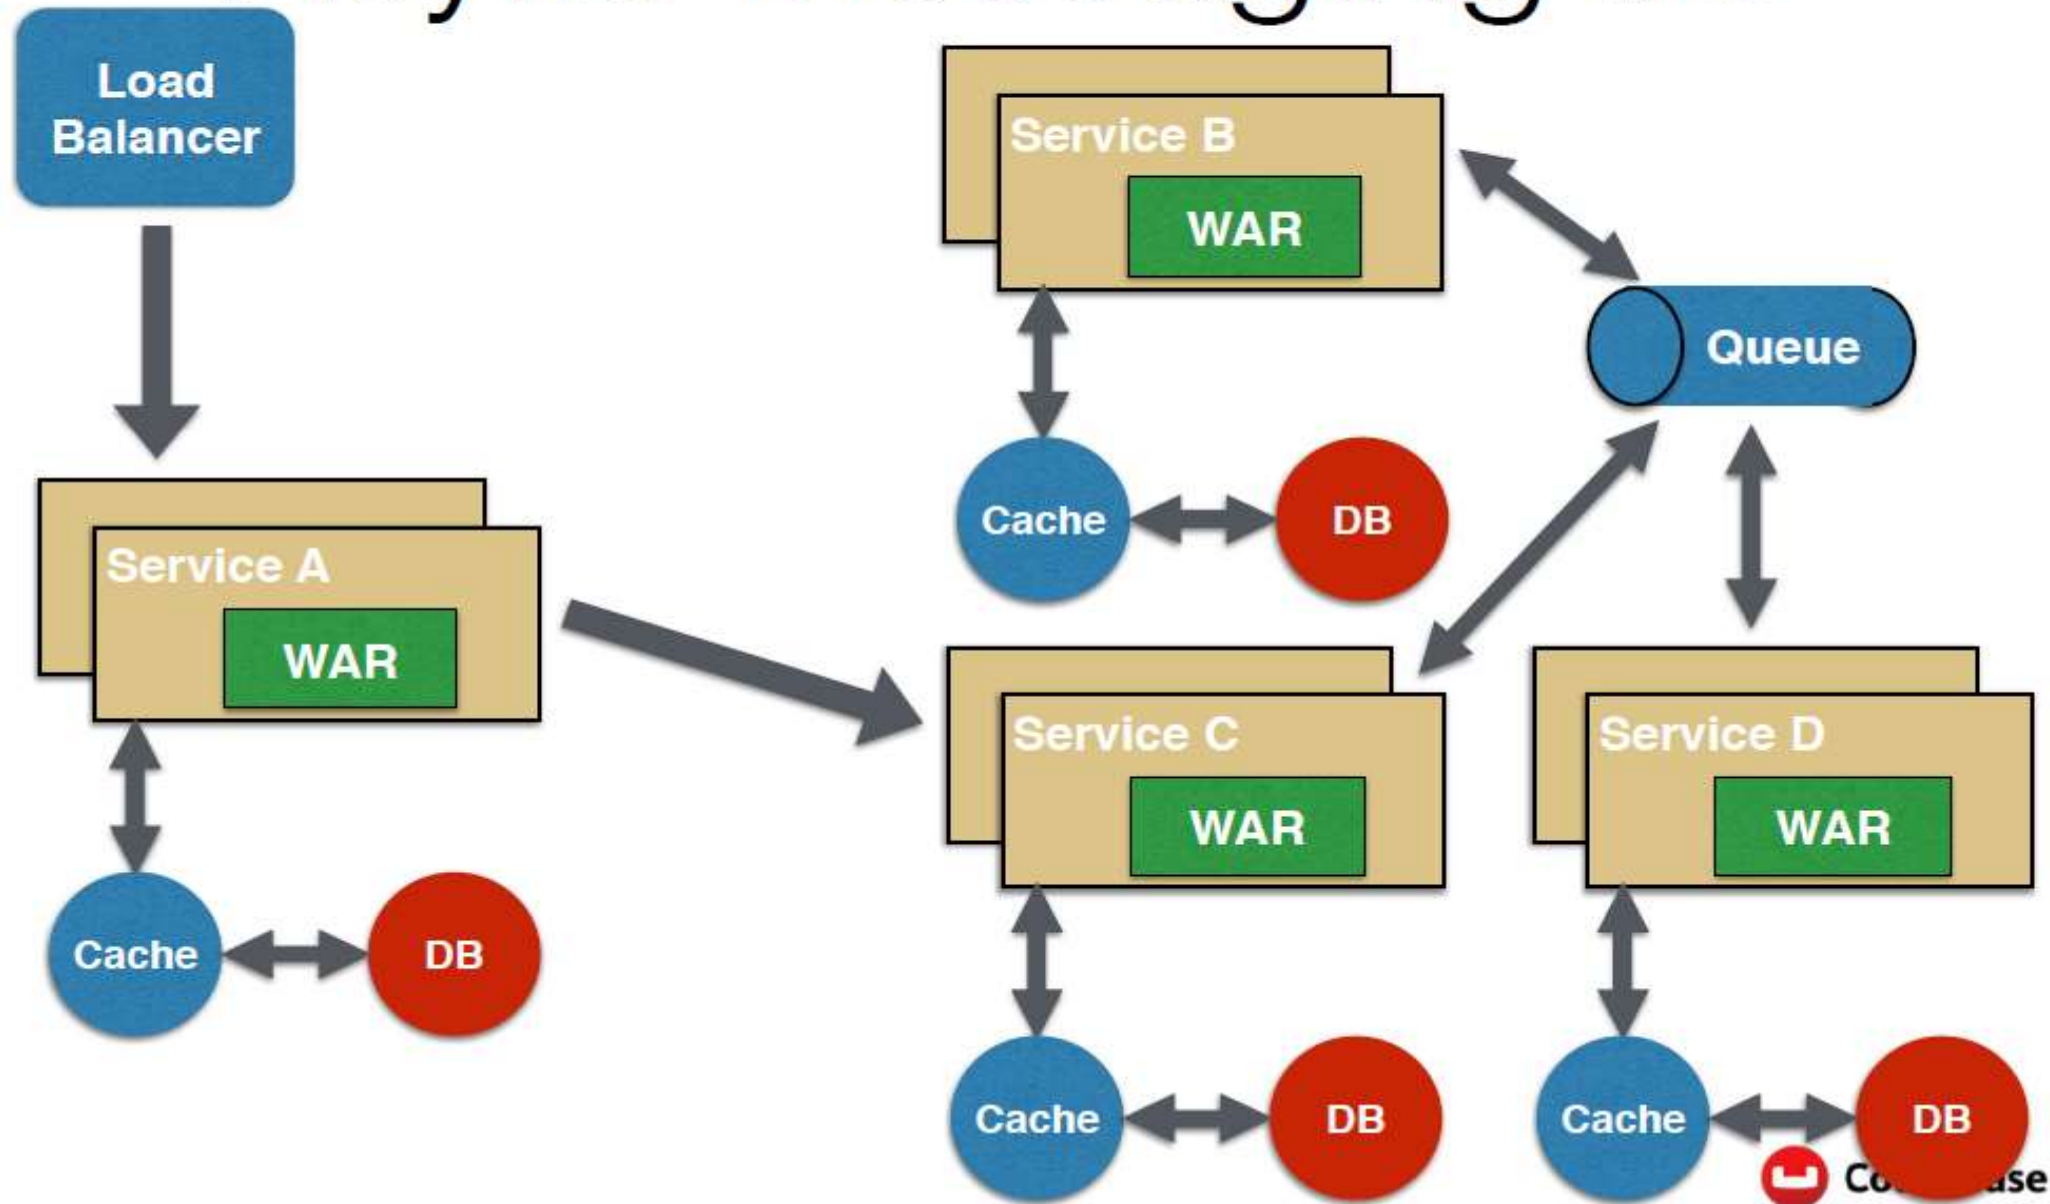
Branch Pattern #4



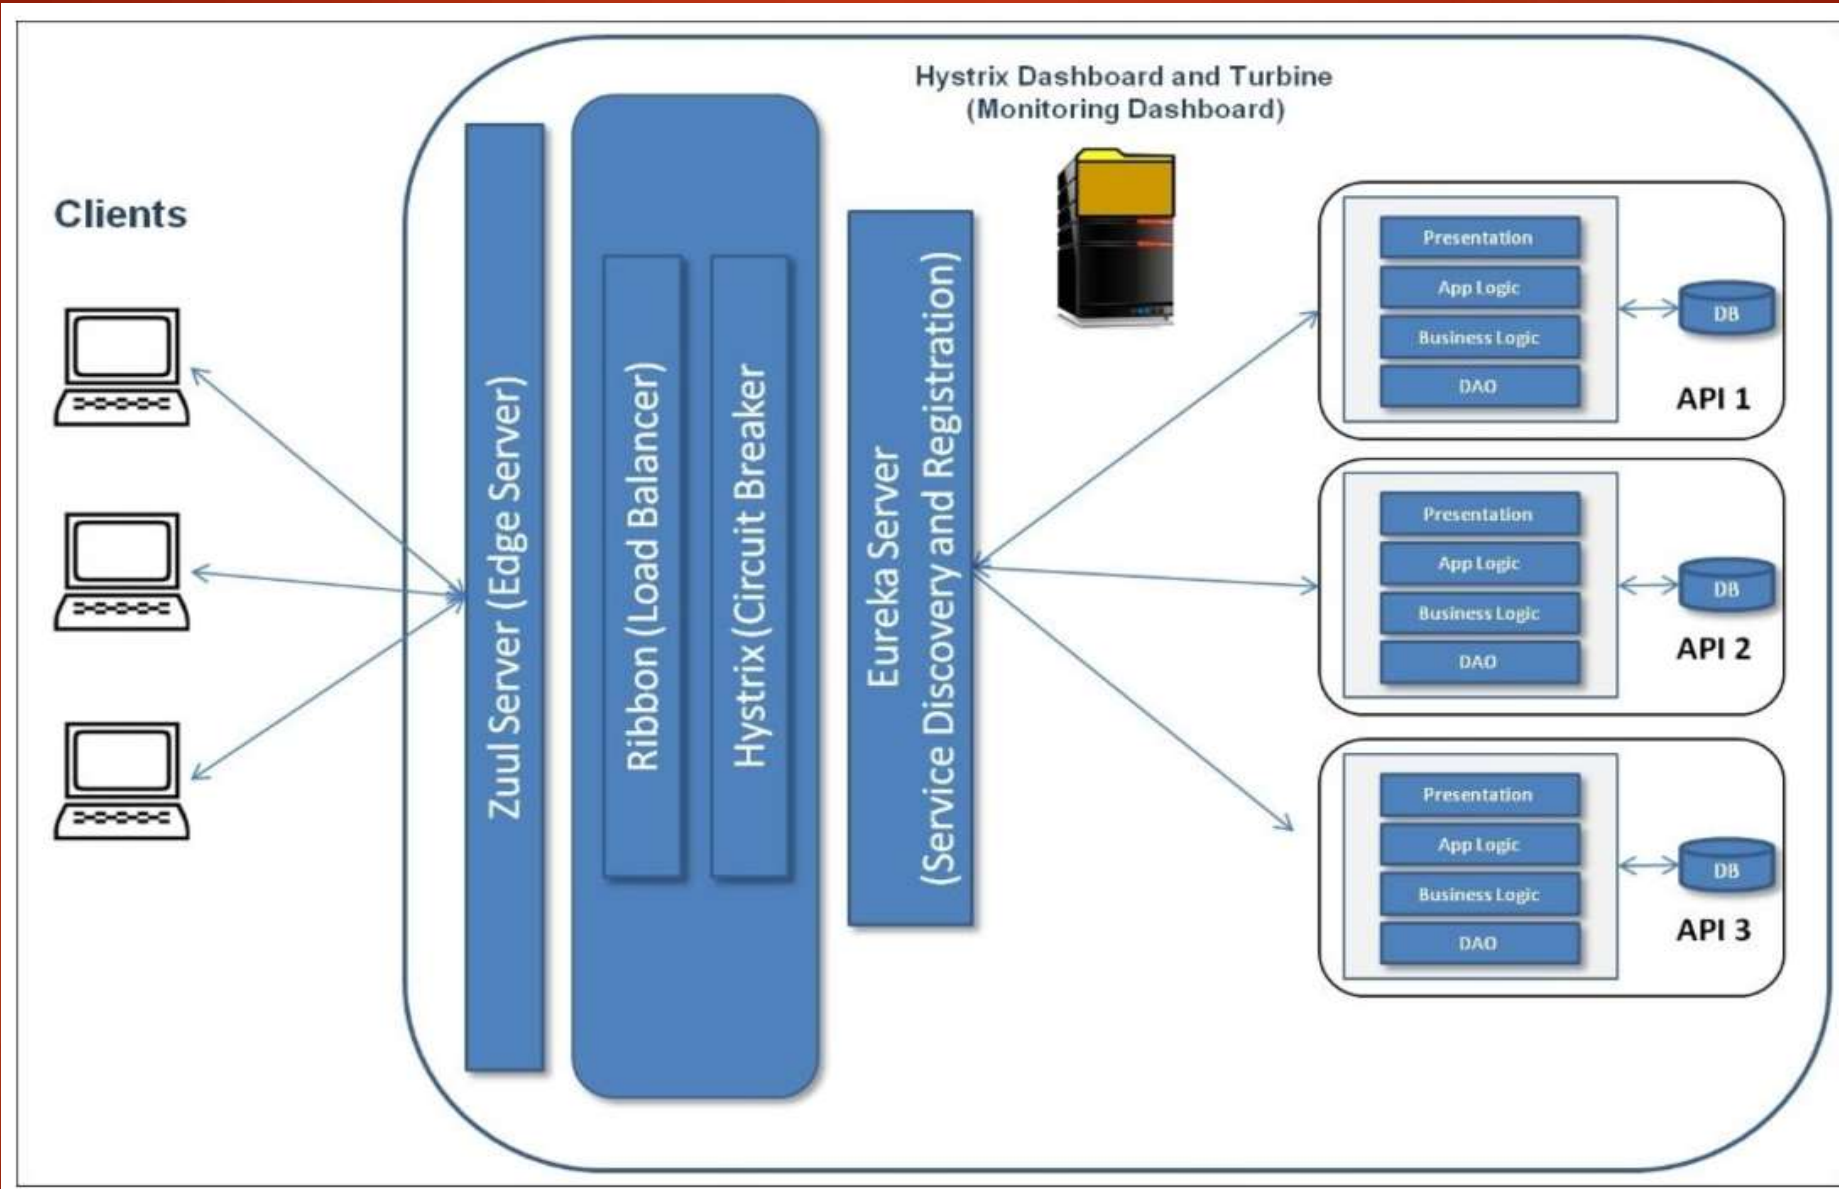
Shared Resources #5



Async Messaging #5



NetflixOSS



localhost:8765/api/account/accounts/customer/2

```
[
  - {
    id: 2,
    customerId: 2,
    number: "222222"
  },
  - {
    id: 6,
    customerId: 2,
    number: "666666"
  },
  - {
    id: 7,
    customerId: 2,
    number: "777777"
  }
]
```

localhost:8765/api/customer/customers/1

```
{
  id: 1,
  pesel: "12345",
  name: "Adam Kowalski",
  type: "INDIVIDUAL",
  - accounts: [
    - {
      id: 1,
      number: "111111"
    },
    - {
      id: 5,
      number: "555555"
    }
  ]
}
```

localhost:2222/accounts/

```
[
  - {
    id: 1,
    customerId: 1,
    number: "111111"
  },
  - {
    id: 2,
    customerId: 2,
    number: "222222"
  },
  - {
    id: 3,
    customerId: 3,
    number: "333333"
  },
  - {
    id: 4,
    customerId: 4,
    number: "444444"
  },
]
```

localhost:3333/customers/

```
[
  - {
    id: 1,
    pesel: "12345",
    name: "Adam Kowalski",
    type: "INDIVIDUAL",
    - accounts: [
      - {
        id: 1,
        number: "111111"
      },
      - {
        id: 5,
        number: "555555"
      }
    ]
  },
  - {
    id: 2,
    pesel: "12346",
    name: "Anna Malinowska",

```

| Name | Value |
|----------------------|-------------|
| total-avail-memory | 487mb |
| environment | test |
| num-of-cpus | 8 |
| current-memory-usage | 149mb (30%) |
| server-uptime | 00:07 |
| registered-replicas | |
| unavailable-replicas | |
| available-replicas | |

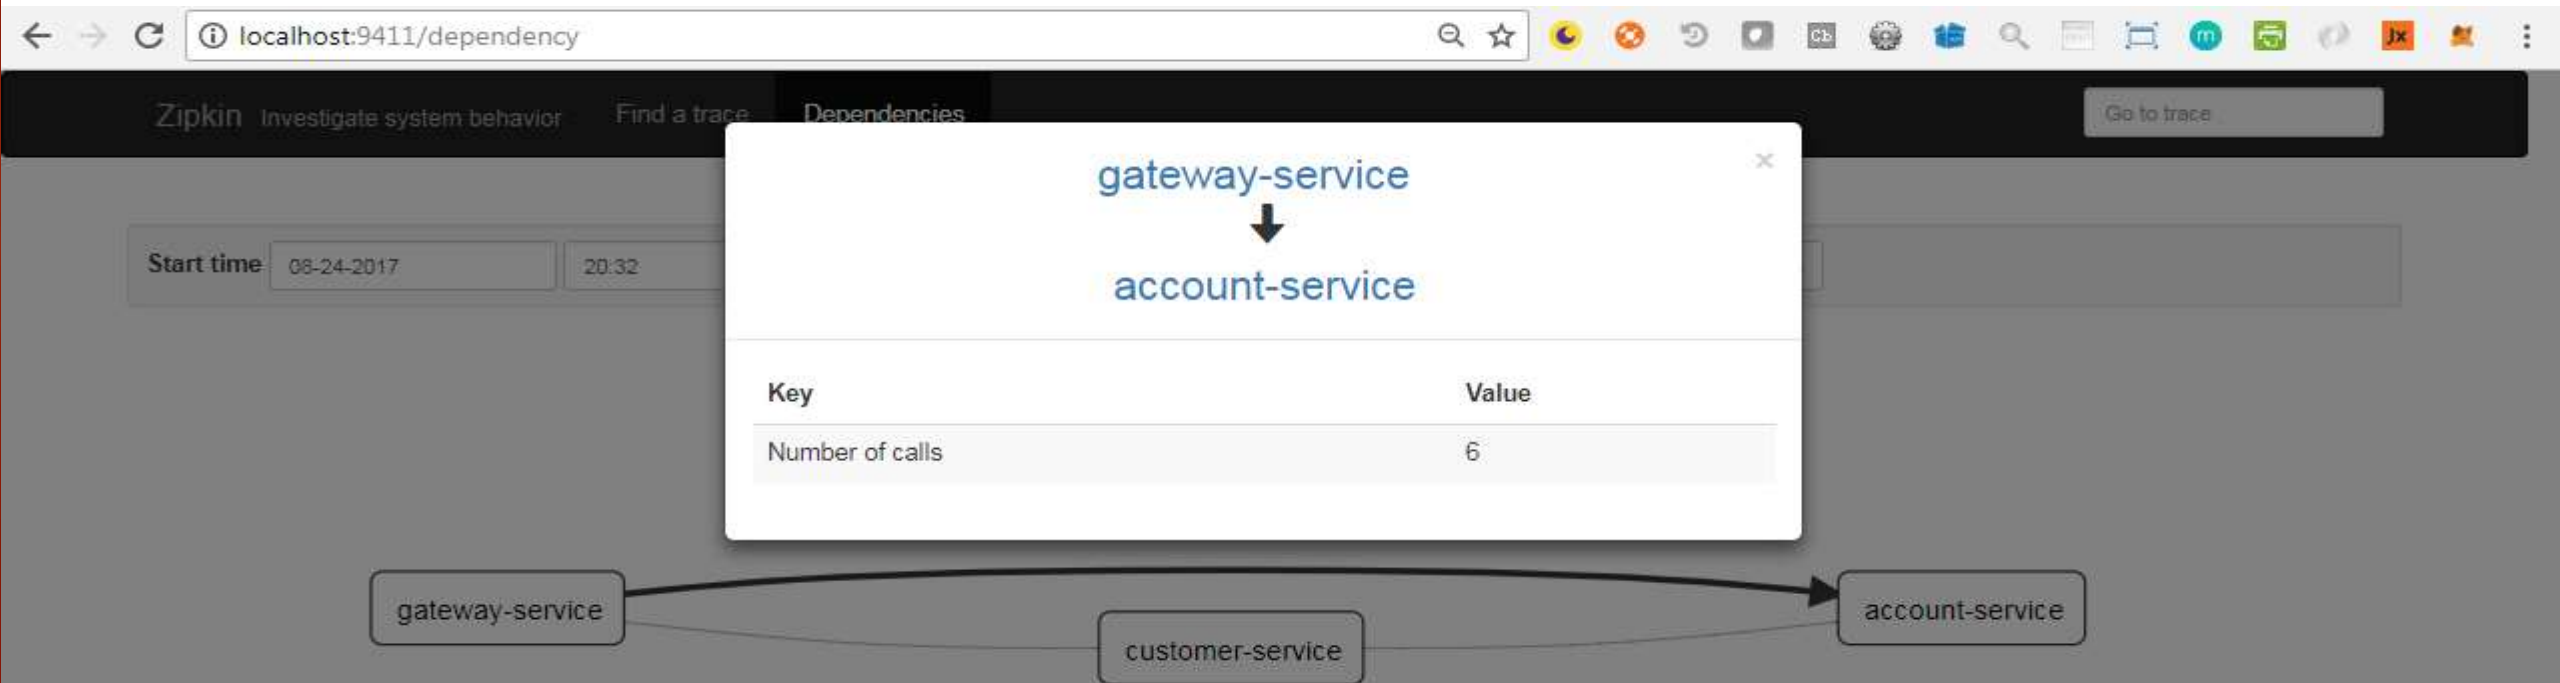
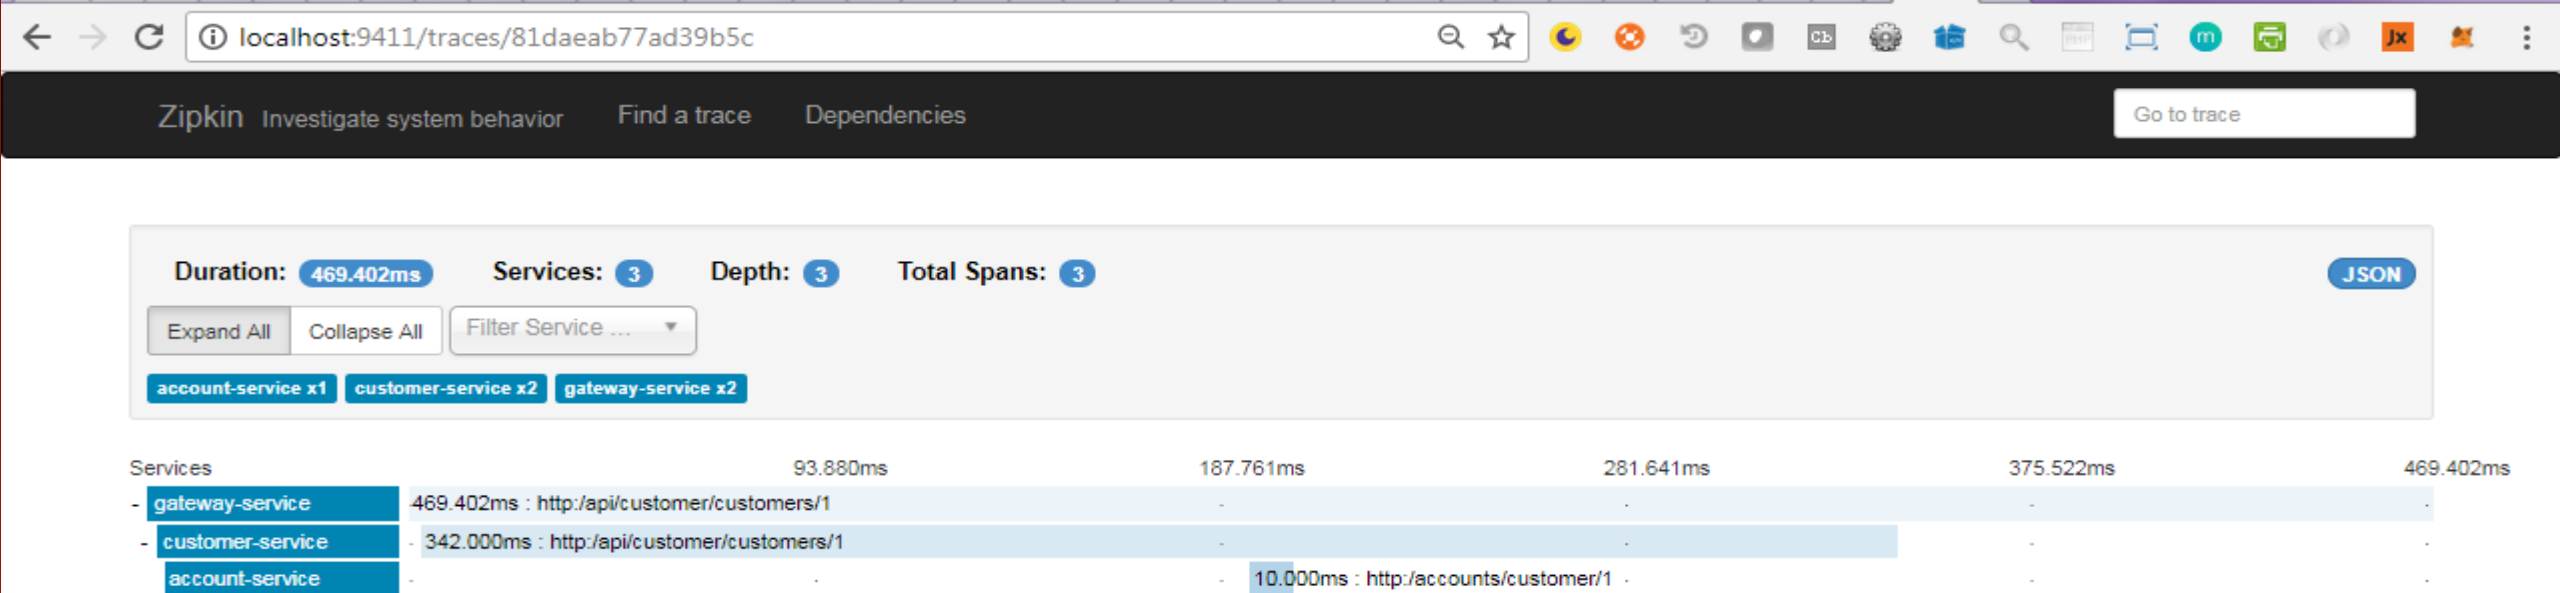
account-service all Start time 08-18-2017 20:21 End time 08-25-2017 20:21

Duration (µs) >= Limit 10 Find Traces

Annotations Query (e.g. "finagle.timeout", "http.path=/foo/bar/ and cluster=foo and cache.miss")

Showing: 4 of 4 Services: account-service Sort: Longest First






| | |
|------------|------------|
| Start time | 08-24-2017 |
|------------|------------|

20:22

| | |
|----------|------------|
| End time | 08-25-2017 |
|----------|------------|

20:25

Analyze Dependencies



| | |
|------------|------------|
| Start time | 08-24-2017 |
|------------|------------|

20:22

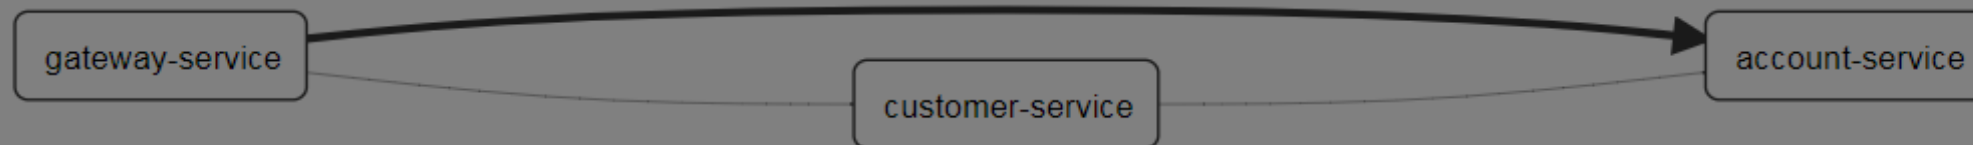
customer-service

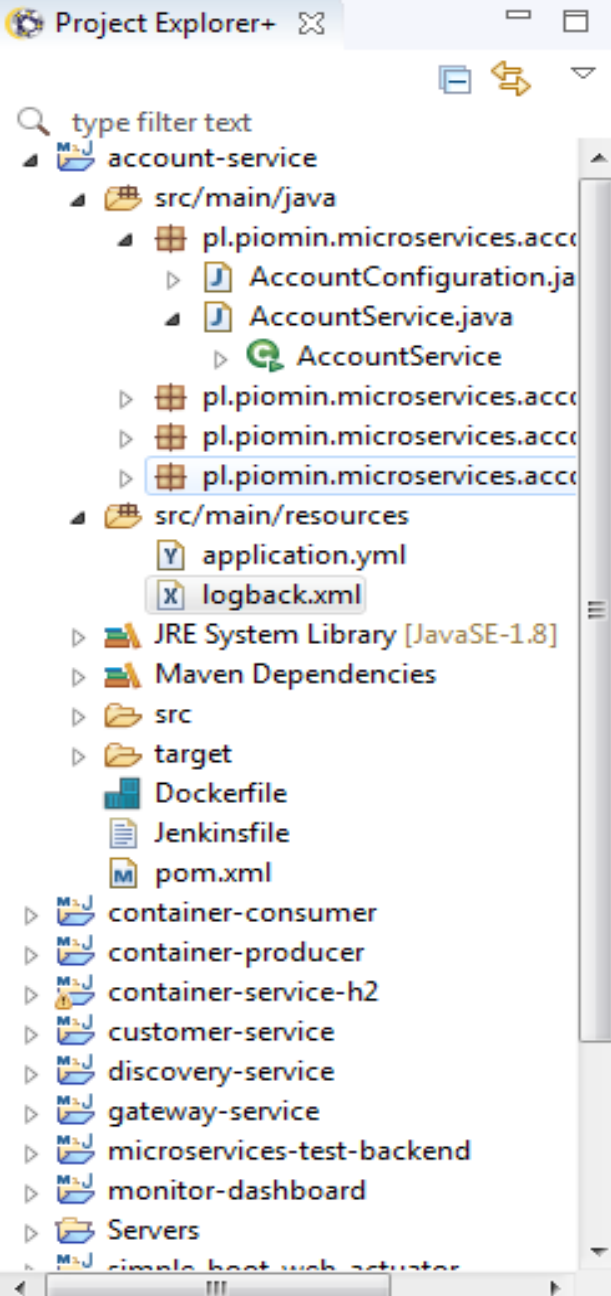
Used by

- gateway-service

Uses

- account-service





```
application.yml ZipkinService.java logback.xml
1 <?xml version="1.0" encoding="UTF-8"?>
2 <configuration>
3
4   <appender name="STDOUT" class="ch.qos.logback.core.ConsoleAppender">
5     <encoder>
6       <pattern>%d{HH:mm:ss.SSS} [%thread, %X{X-B3-TraceId:-},%X{X-B3-SpanId:-}] %-5level %logger{36} %n
7     </encoder>
8   </appender>
9
10  <appender name="STASH"
11    class="net.logstash.logback.appender.LogstashTcpSocketAppender">
12    <destination>127.0.0.1:9000</destination>
13
14    <encoder
15      class="net.logstash.logback.encoder.LoggingEventCompositeJsonEncoder">
16        <providers>
17          <mdc /> <!-- MDC variables on the Thread will be written as JSON fields -->
18          <context /> <!-- Outputs entries from logback's context -->
19          <version /> <!-- Logstash json format version, the @version field in the output -->
20          <logLevel />
21          <loggerName />
22
23          <pattern>
24            <pattern>
25              {
26                "serviceName": "account-service"
27              }
28            </pattern>
29          </pattern>
30
```

Design Source

Terminal+ Console Markers Git Staging

DiscoveryService [Java Application] C:\Program Files\Java\jre1.8.0_45\bin\javaw.exe (Aug 25, 2017, 7:55:02 PM)

```
2017-08-25 19:55:09.103 INFO 8608 --- [Thread-11] o.s.c.n.e.server.EurekaServerBootstrap : Initialize
2017-08-25 19:55:09.103 INFO 8608 --- [Thread-11] c.n.e.r.PeerAwareInstanceRegistryImpl : Got 1 inst
2017-08-25 19:55:09.103 INFO 8608 --- [Thread-11] c.n.e.r.PeerAwareInstanceRegistryImpl : Renew thre
2017-08-25 19:55:09.103 INFO 8608 --- [Thread-11] c.n.e.r.PeerAwareInstanceRegistryImpl : Changing s
2017-08-25 19:55:09.108 INFO 8608 --- [Thread-11] e.s.EurekaServerInitializerConfiguration : Started Eu
2017-08-25 19:55:09.237 INFO 8608 --- [main] s.b.c.e.t.TomcatEmbeddedServletContainer : Tomcat s
```

logstash-simple.conf

input {

tcp {

port => 9000 codec => "json"

}

}

output {

elasticsearch {

hosts => ["localhost:9200"] index =>

"micro-%{serviceName}"

}

}

```

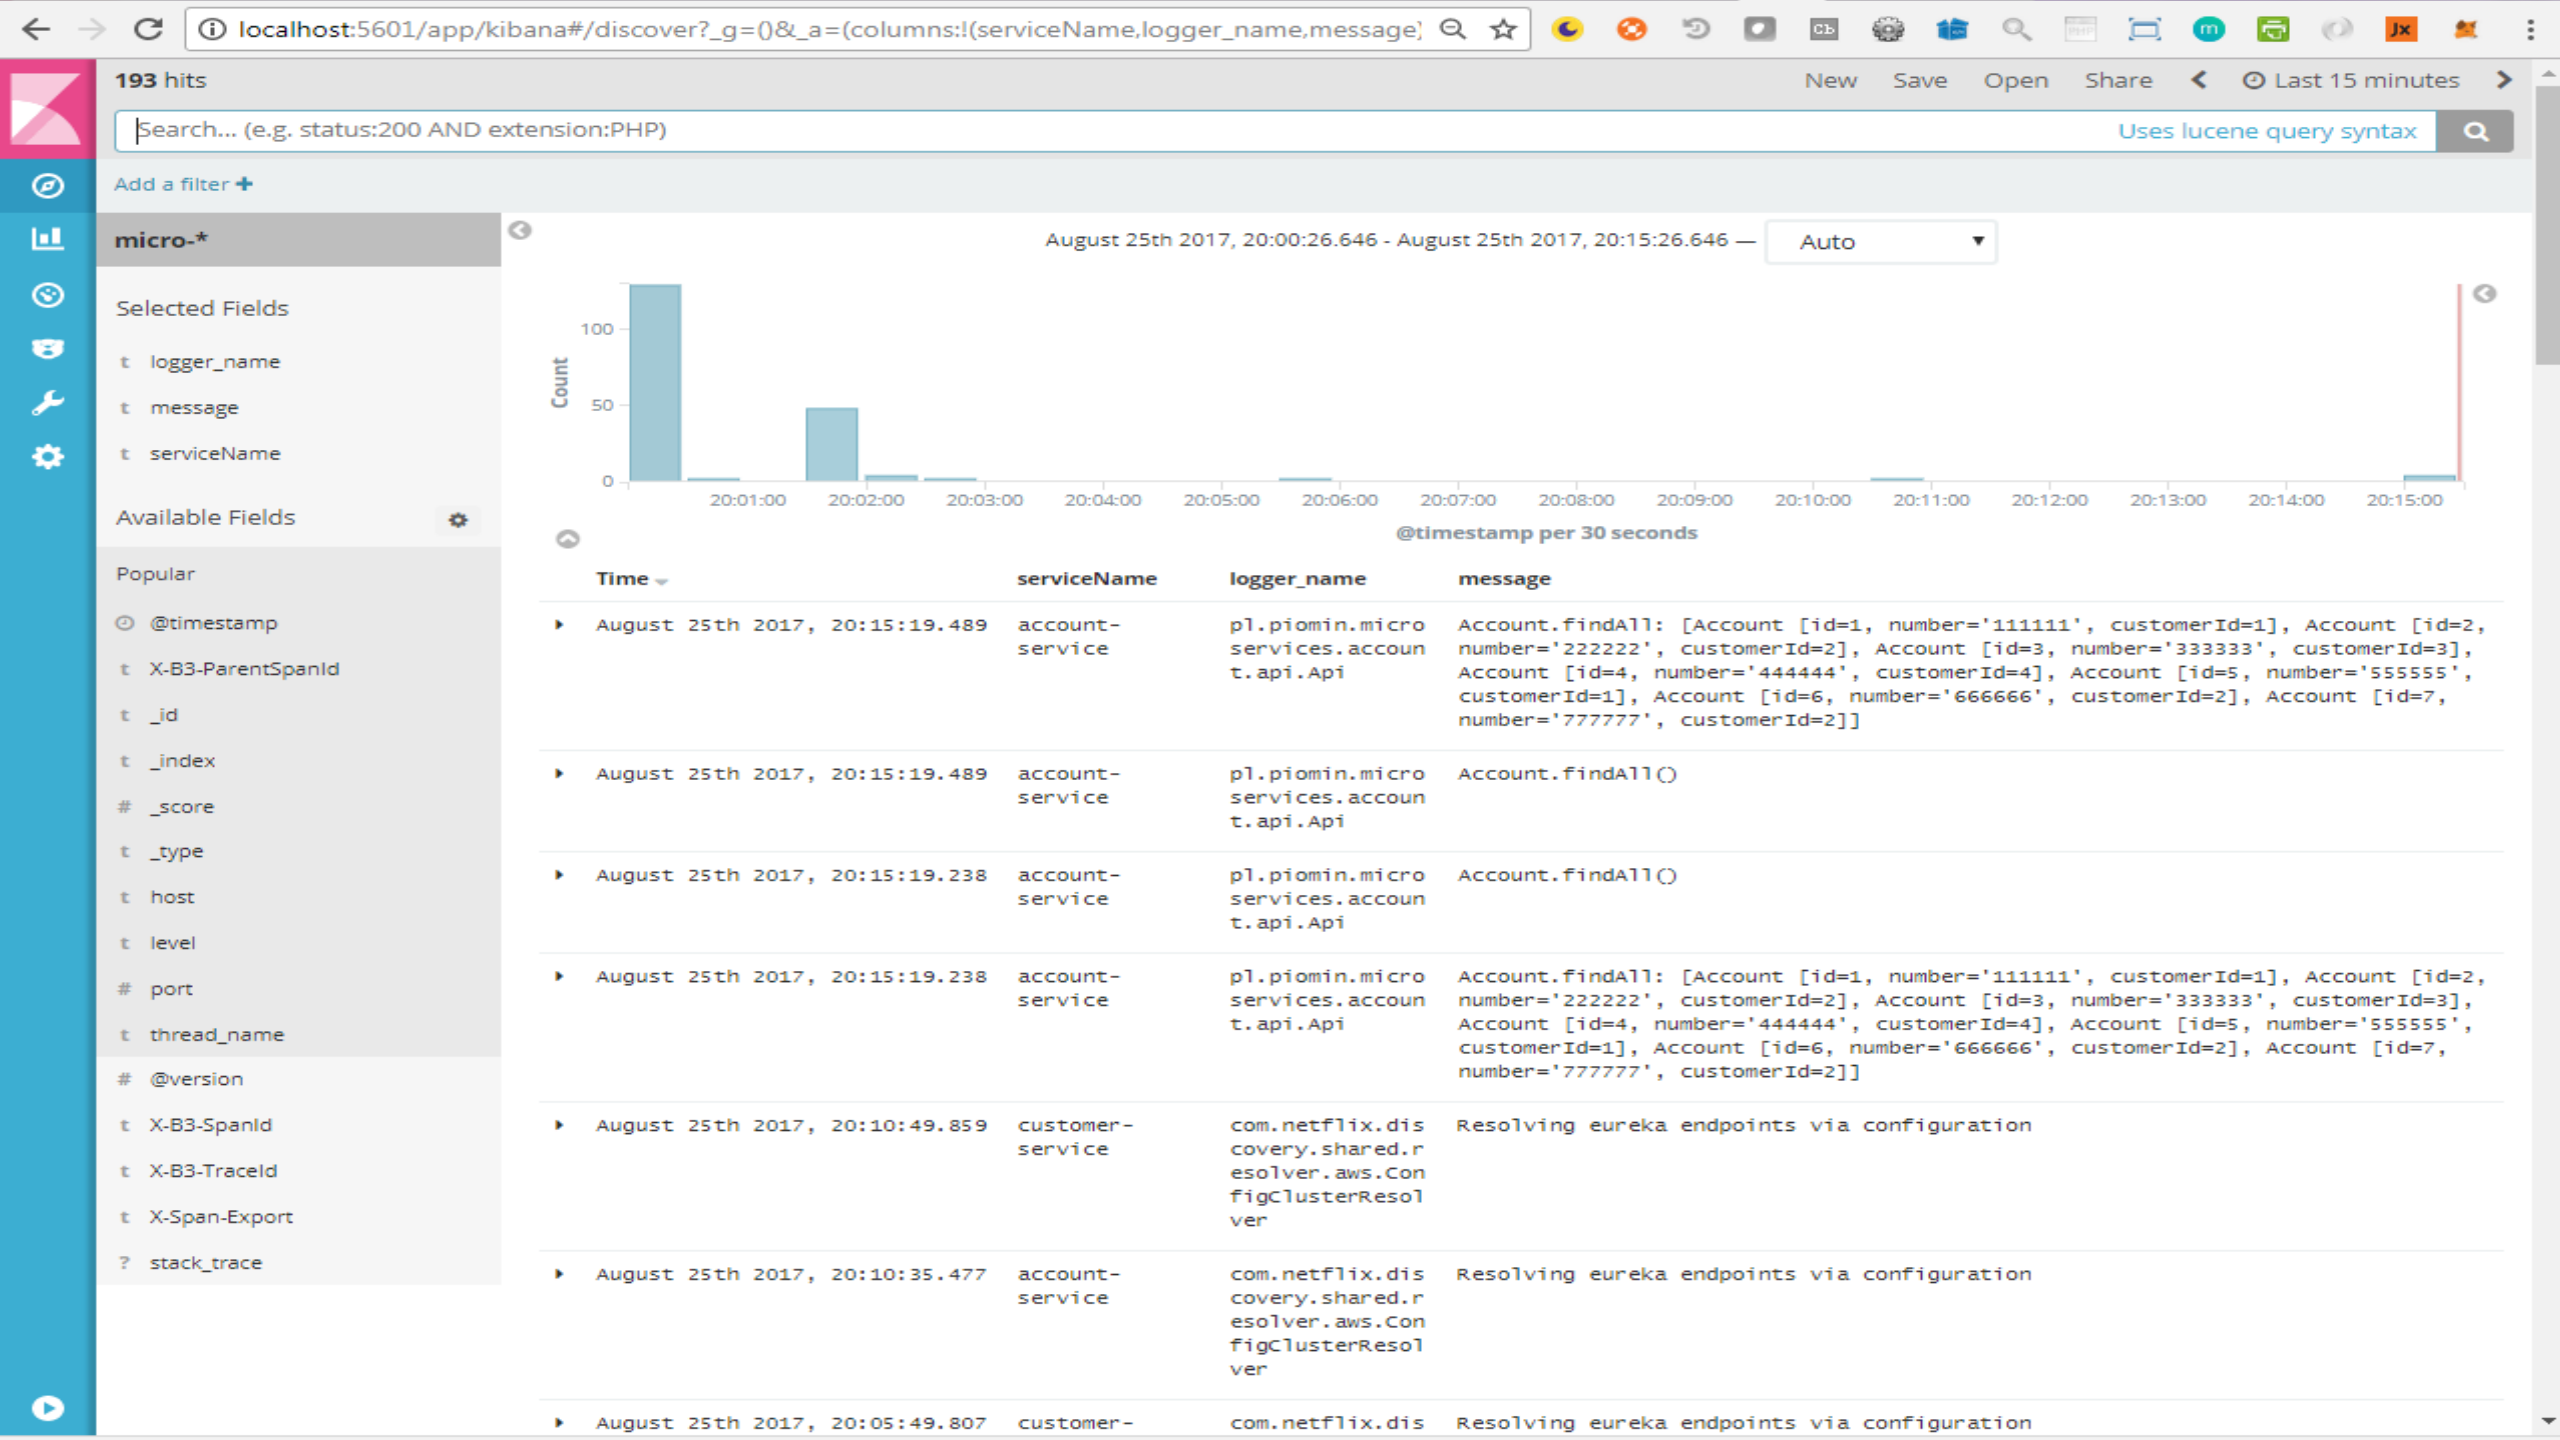
Elasticsearch 5.5.0
[2017-08-25T19:59:09,831][INFO ][o.e.p.PluginsService] [bk8FBjQ] loaded module [transport-netty3]
[2017-08-25T19:59:09,831][INFO ][o.e.p.PluginsService] [bk8FBjQ] loaded module [transport-netty4]
[2017-08-25T19:59:09,832][INFO ][o.e.p.PluginsService] [bk8FBjQ] no plugins loaded
[2017-08-25T19:59:13,873][INFO ][o.e.d.DiscoveryModule] [bk8FBjQ] using discovery type [zen]
[2017-08-25T19:59:14,849][INFO ][o.e.n.Node] [bk8FBjQ] initialized
[2017-08-25T19:59:14,850][INFO ][o.e.n.Node] [bk8FBjQ] starting
[2017-08-25T19:59:15,649][INFO ][o.e.t.TransportService] [bk8FBjQ] publish_address {127.0.0.1:9300}, bound_addresses {127.0.0.1:9300}, {:::1:9300}
[2017-08-25T19:59:18,810][INFO ][o.e.c.s.ClusterService] [bk8FBjQ] new_master {bk8FBjQ}<bk8FBjQ0SYa5UheIdXWSmA><37-NBsDhT9eM6tFGItmxwA><127.0.0.1><127.0.0.1:9300>, reason: zen-disco-elected-as-master (<0> nodes joined)
[2017-08-25T19:59:19,102][INFO ][o.e.h.n.Netty4HttpServerTransport] [bk8FBjQ] publish_address {127.0.0.1:9200}, bound_addresses {127.0.0.1:9200}, {:::1:9200}
[2017-08-25T19:59:19,103][INFO ][o.e.n.Node] [bk8FBjQ] started
[2017-08-25T19:59:19,301][INFO ][o.e.g.GatewayService] [bk8FBjQ] recovered [3] indices into cluster_state
[2017-08-25T19:59:24,095][INFO ][o.e.c.r.a.AllocationService] [bk8FBjQ] Cluster health status changed from [RED] to [YELLOW] (reason: [shards started [[micro-account-service] [2]] ...]).

```

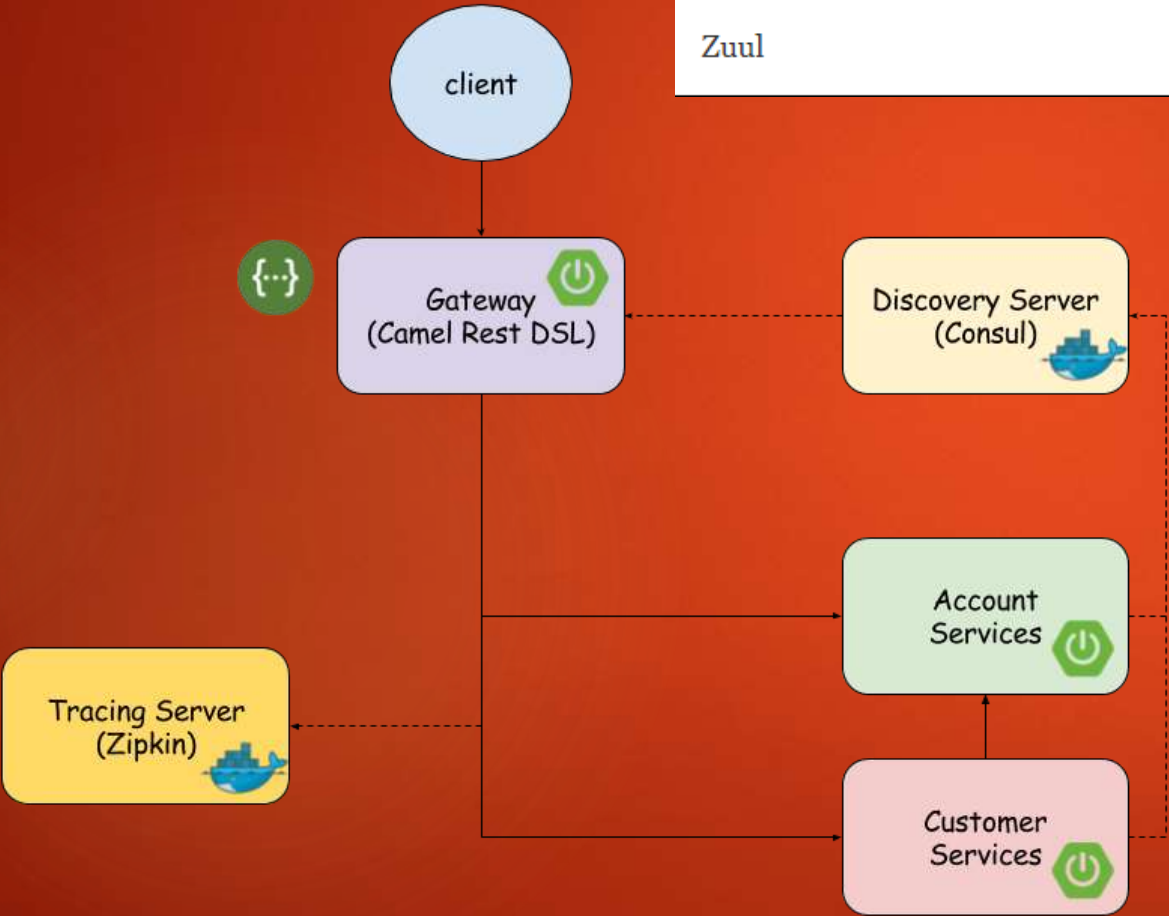
```

Kibana Server
log [14:30:20.490] [info][status][plugin:kibana@5.5.0] Status changed from uninitialized to green - Ready
log [14:30:20.585] [info][status][plugin:elasticsearch@5.5.0] Status changed from uninitialized to yellow - Waiting for Elasticsearch
log [14:30:20.608] [info][status][plugin:console@5.5.0] Status changed from uninitialized to green - Ready
log [14:30:20.630] [info][status][plugin:metrics@5.5.0] Status changed from uninitialized to green - Ready
log [14:30:21.245] [info][status][plugin:timelion@5.5.0] Status changed from uninitialized to green - Ready
log [14:30:21.252] [info][listening] Server running at http://localhost:5601
log [14:30:21.254] [info][status][ui settings] Status changed from uninitialized to yellow - Elasticsearch plugin is yellow
log [14:30:21.506] [info][status][plugin:elasticsearch@5.5.0] Status changed from yellow to green - Kibana index ready
log [14:30:21.509] [info][status][ui settings] Status changed from yellow to green - Ready

```



| Netflix Component Name | Functionality |
|------------------------|------------------------------------|
| Eureka | Service Registration and Discovery |
| Ribbon | Dynamic Routing and Load Balancer |
| Hystrix | Circuit Breaker |
| Zuul | Edge Server |



Characteristics



PostgreSQL



mongoDB



Scala



redis



Apache TomEE



Thank You

